

# Microservice Application Architecture...



# Day 1 ReCap

- Monolithic Overview, Pros & Cons
- Microservices Overview, Pros & Cons
- The 12 Factor Apps
- Microservices Challenges
- Building Blocks of Microservices
- Docker Architecture
- Docker Compose
- Contracts, Cohesion, Coupling, Observability, Bounded Contexts
- Service Decomposition
- Microservice Boundaries
- Microservices Communication
- Message Patterns - Pub/Sub, Anycast, Unicast, Multicast
- Message Brokers

# APACHE KAFKA



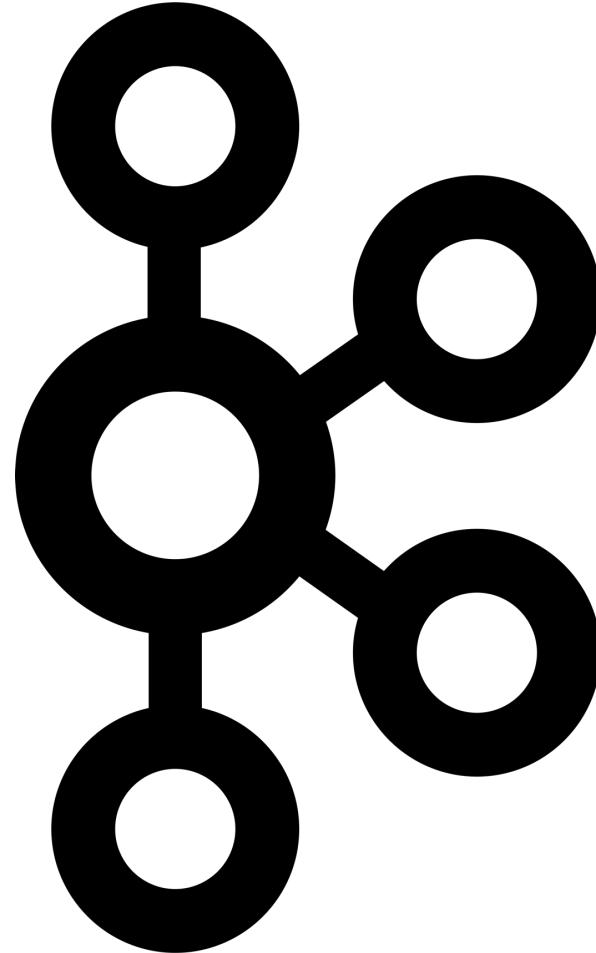
Kafka® is used for building real-time data pipelines and streaming apps.

# APACHE KAFKA



**It is horizontally scalable, fault-tolerant, wicked fast, and runs in production in thousands of companies.**

# APACHE KAFKA



## PUBLISH & SUBSCRIBE

- Read and write streams of data like a messaging system.

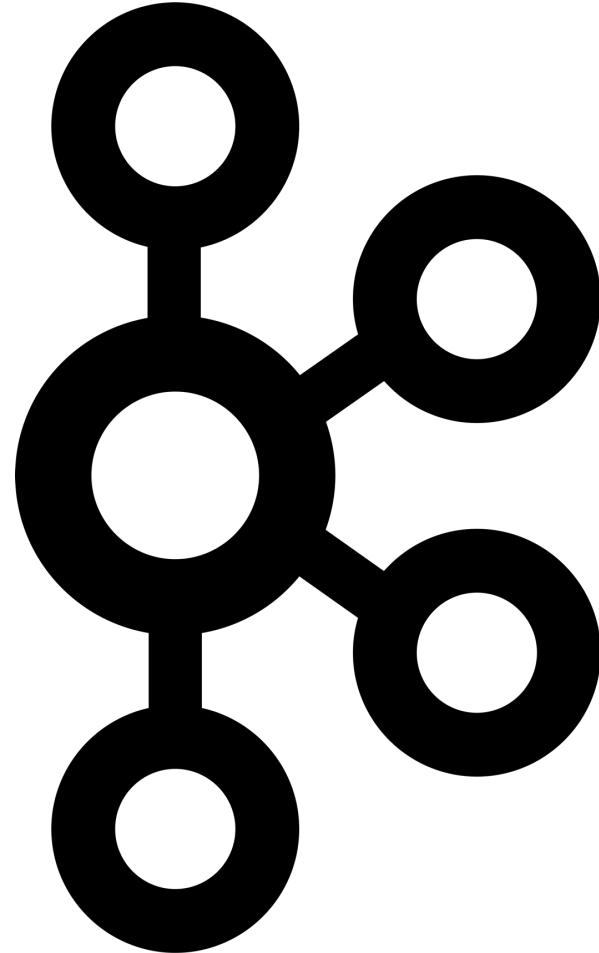
## PROCESS

- Write scalable stream processing applications that react to events in real-time.

## STORE

- Store streams of data safely in a distributed, replicated, fault-tolerant cluster.

# APACHE KAFKA



- **Messages are stored in flat files**
- **Consumers able to ask messages that are based on an offset**
- **Routing capable through Zookeeper**
- **Rules are not available**
- **Stats available through JMX**
- **Great performance and low resource usage**

# ZEROMQ \ZERO-EM-QUEUE\:



- Ø Connect your code in any language, on any platform.
- Ø Carries messages across inproc, IPC, TCP, TIPC, multicast.
- Ø Smart patterns like pub-sub, push-pull, and router-dealer.
- Ø High-speed asynchronous I/O engines, in a tiny library.
- Ø Backed by a large and active open source community.
- Ø Supports every modern language and platform.
- Ø Build any architecture: centralized, distributed, small, or large.

# APACHE ACTIVEMQ



Apache ActiveMQ™ is the most popular and powerful open source messaging and Integration Patterns server.

# APACHE ACTIVEMQ



**Supports a variety of Cross Language Clients and Protocols including:**

- OpenWire, Stomp, AMQP, MQTT

**Full support for the Enterprise Integration Patterns**

- JMS client and the Message Broker

# **MONITORING, API SECURITY AND INIT CONTAINERS**

# MONITORING

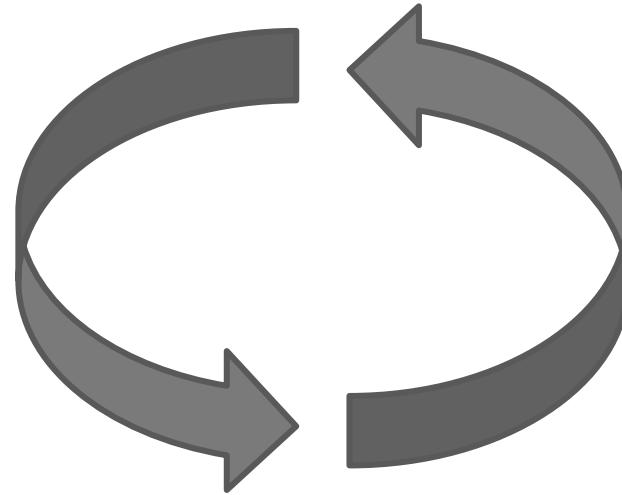
Why monitor?



# MONITORING



Quality  
Assurance



Continuous  
Improvement

# MONITORING

Not about collecting data

- Why vs How

# **MONITORING**

## **Why monitor?**

**Know when things go wrong**

**Debug and learn from issues**

**Historical view/Trends**

**Drive technical/business decisions**

**Feed into other systems/processes (QA, Security, automation)**

# **MONITORING**

## **Challenges**

**Monitoring tools are limited**

- Technically
- Conceptually

**Tools don't scale well and are difficult to manage**

**Operational practices don't align with business needs**

# MONITORING

## Challenges

**Example:**

**Customers care about increased latency and it's in your SLAs. Monitoring only supports alerting on individual machines performance and not the entire clusters.**

**Result: Engineers get a bunch of non critical/false-positive alerts and are woken up for non-issues. This leads to fatigue and ignoring alerts.**

# MONITORING



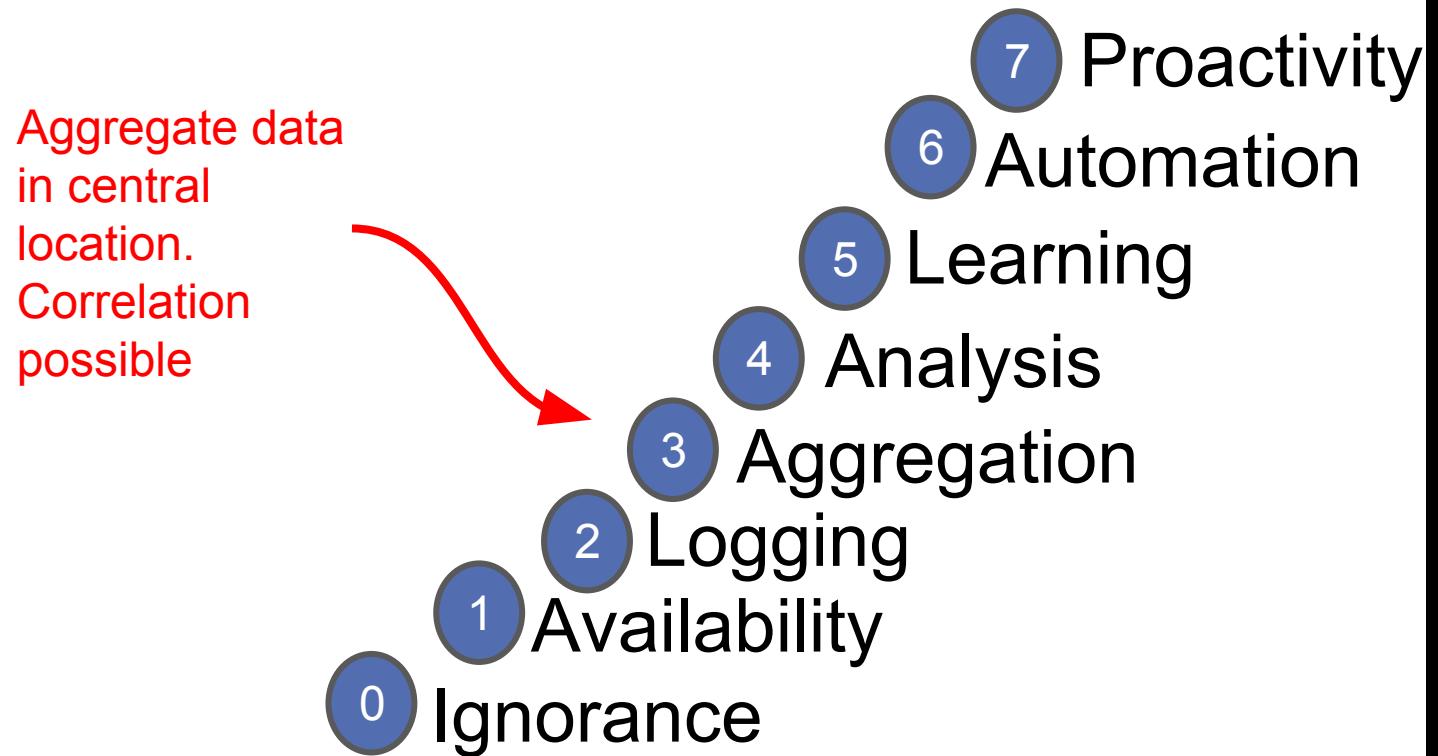
# MONITORING



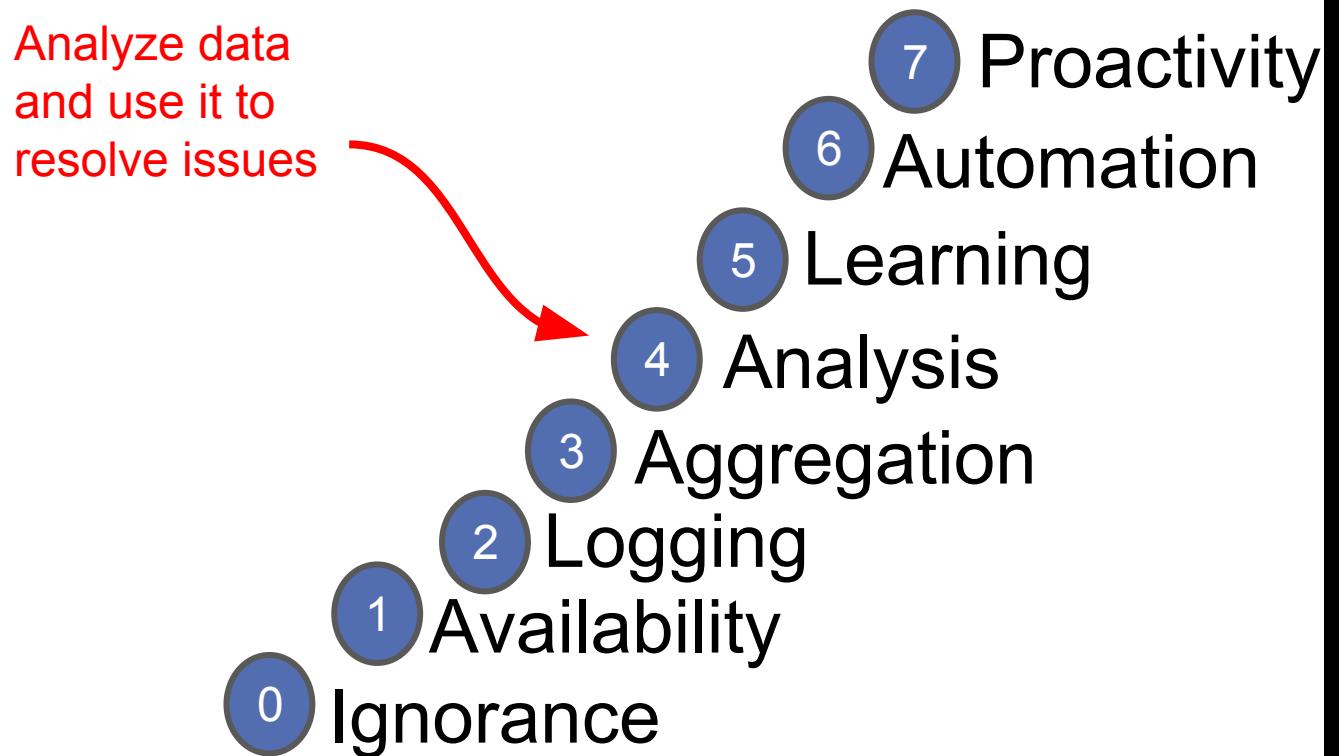
# MONITORING



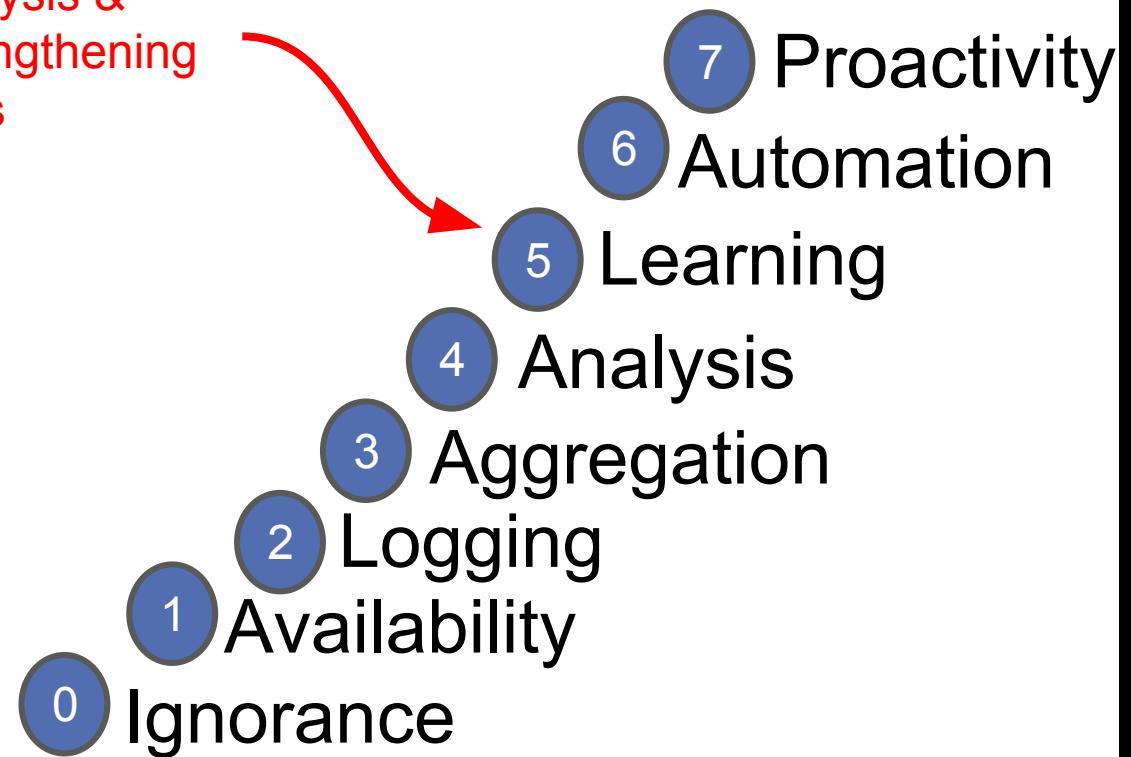
# MONITORING



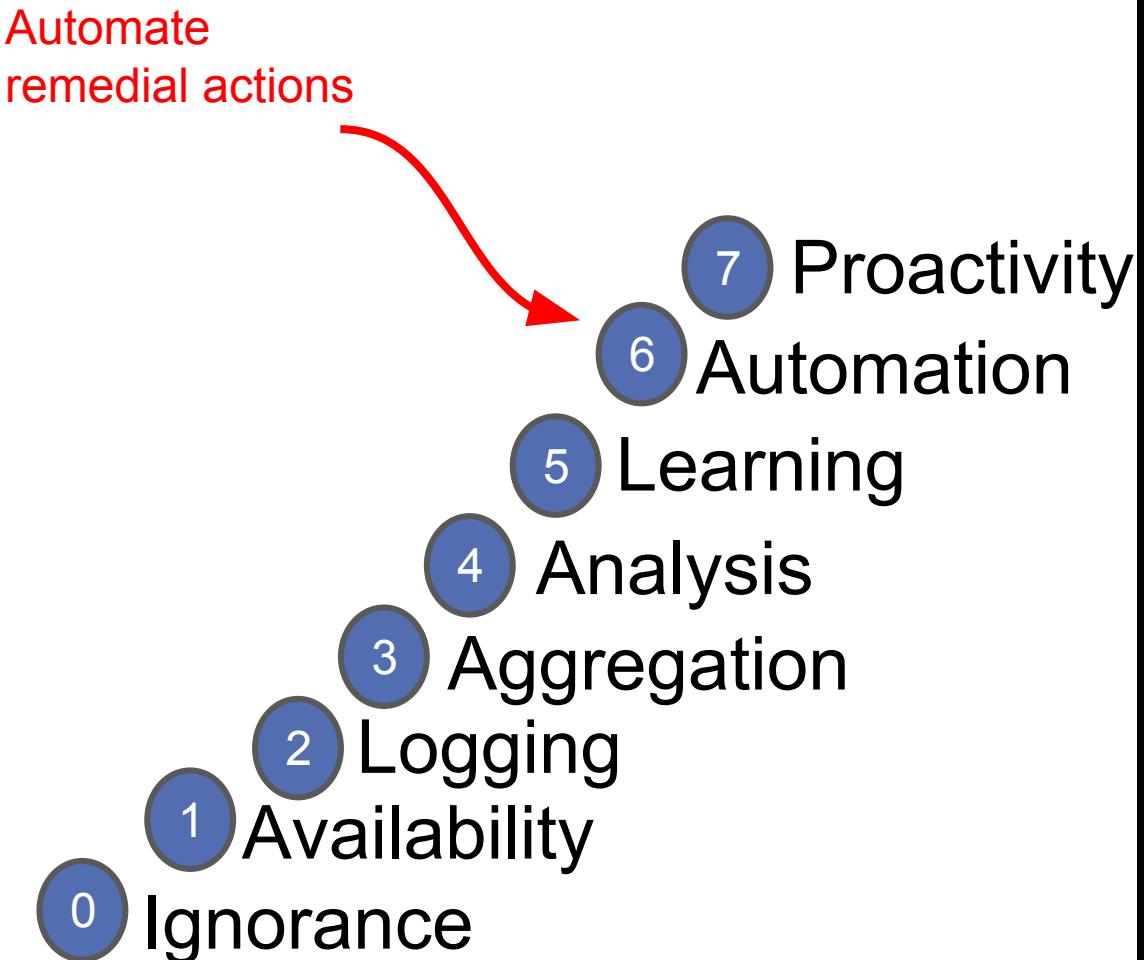
# MONITORING



# MONITORING



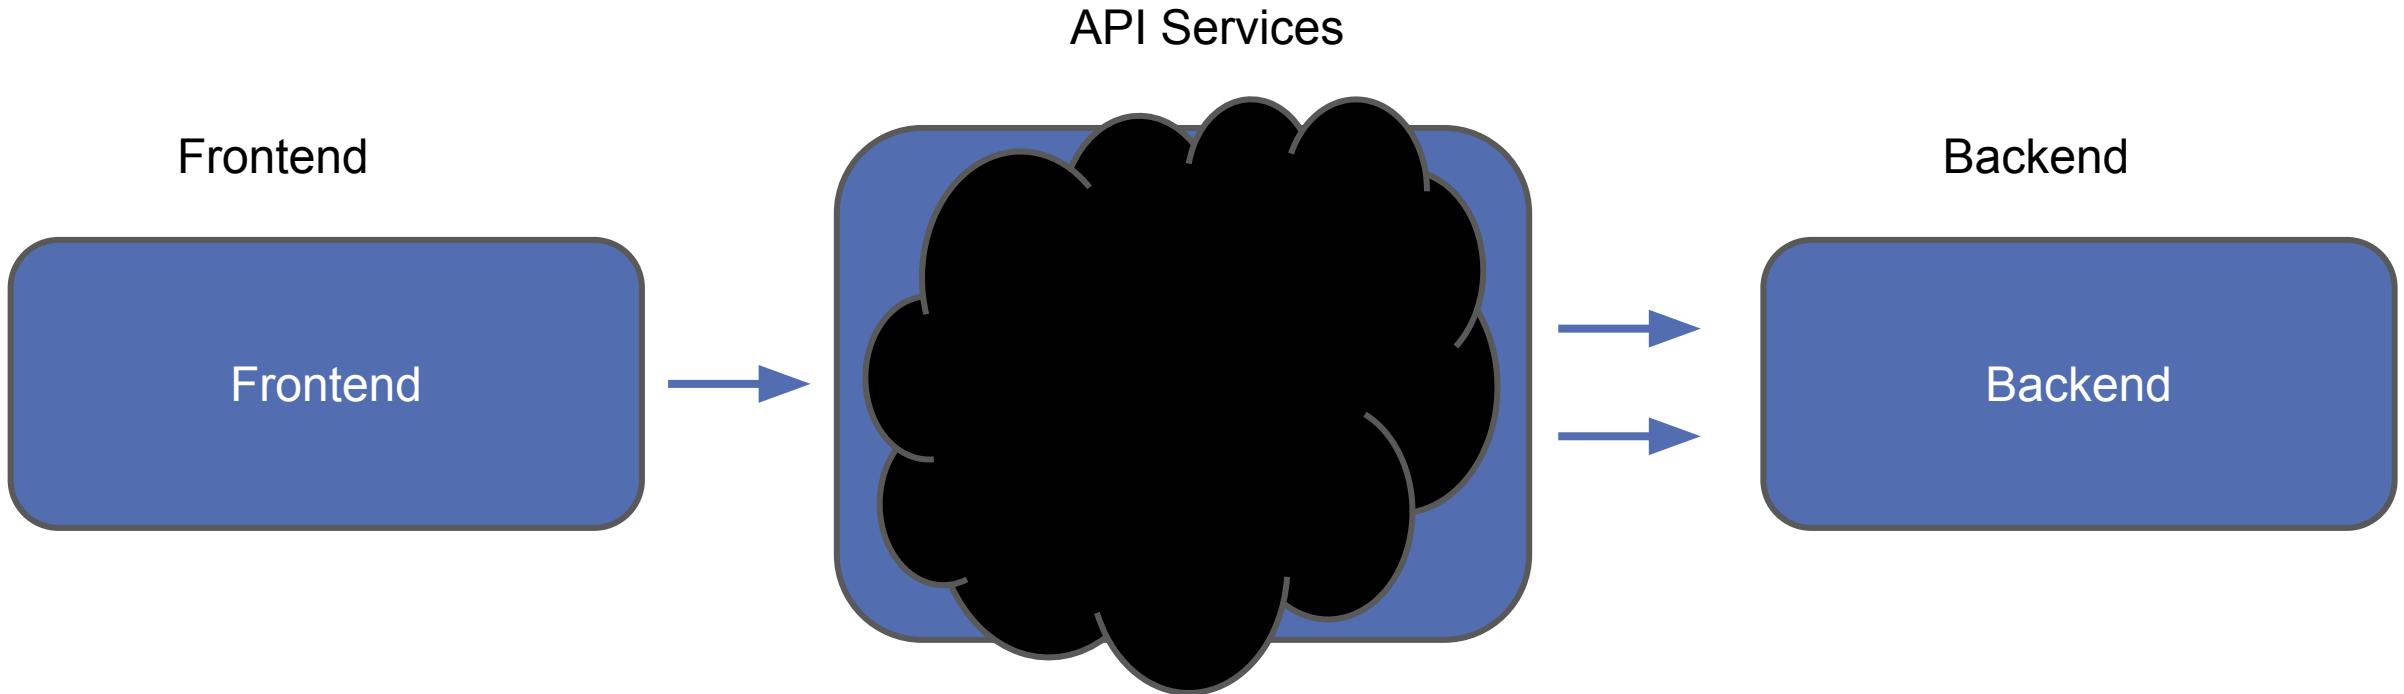
# MONITORING



# MONITORING

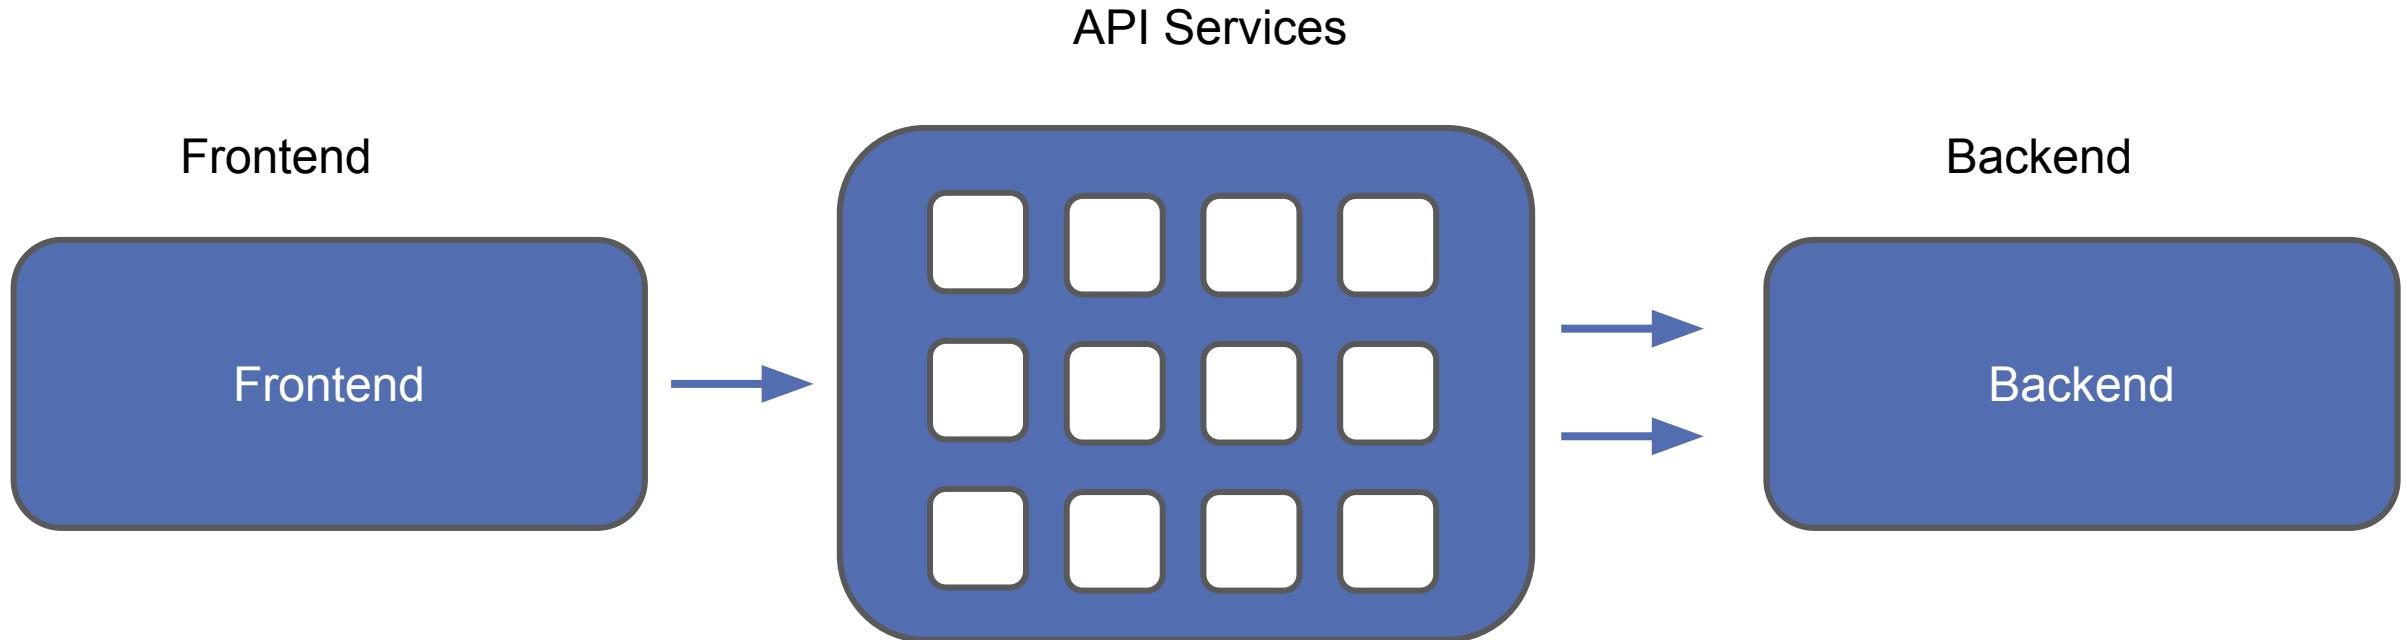


# MONITORING



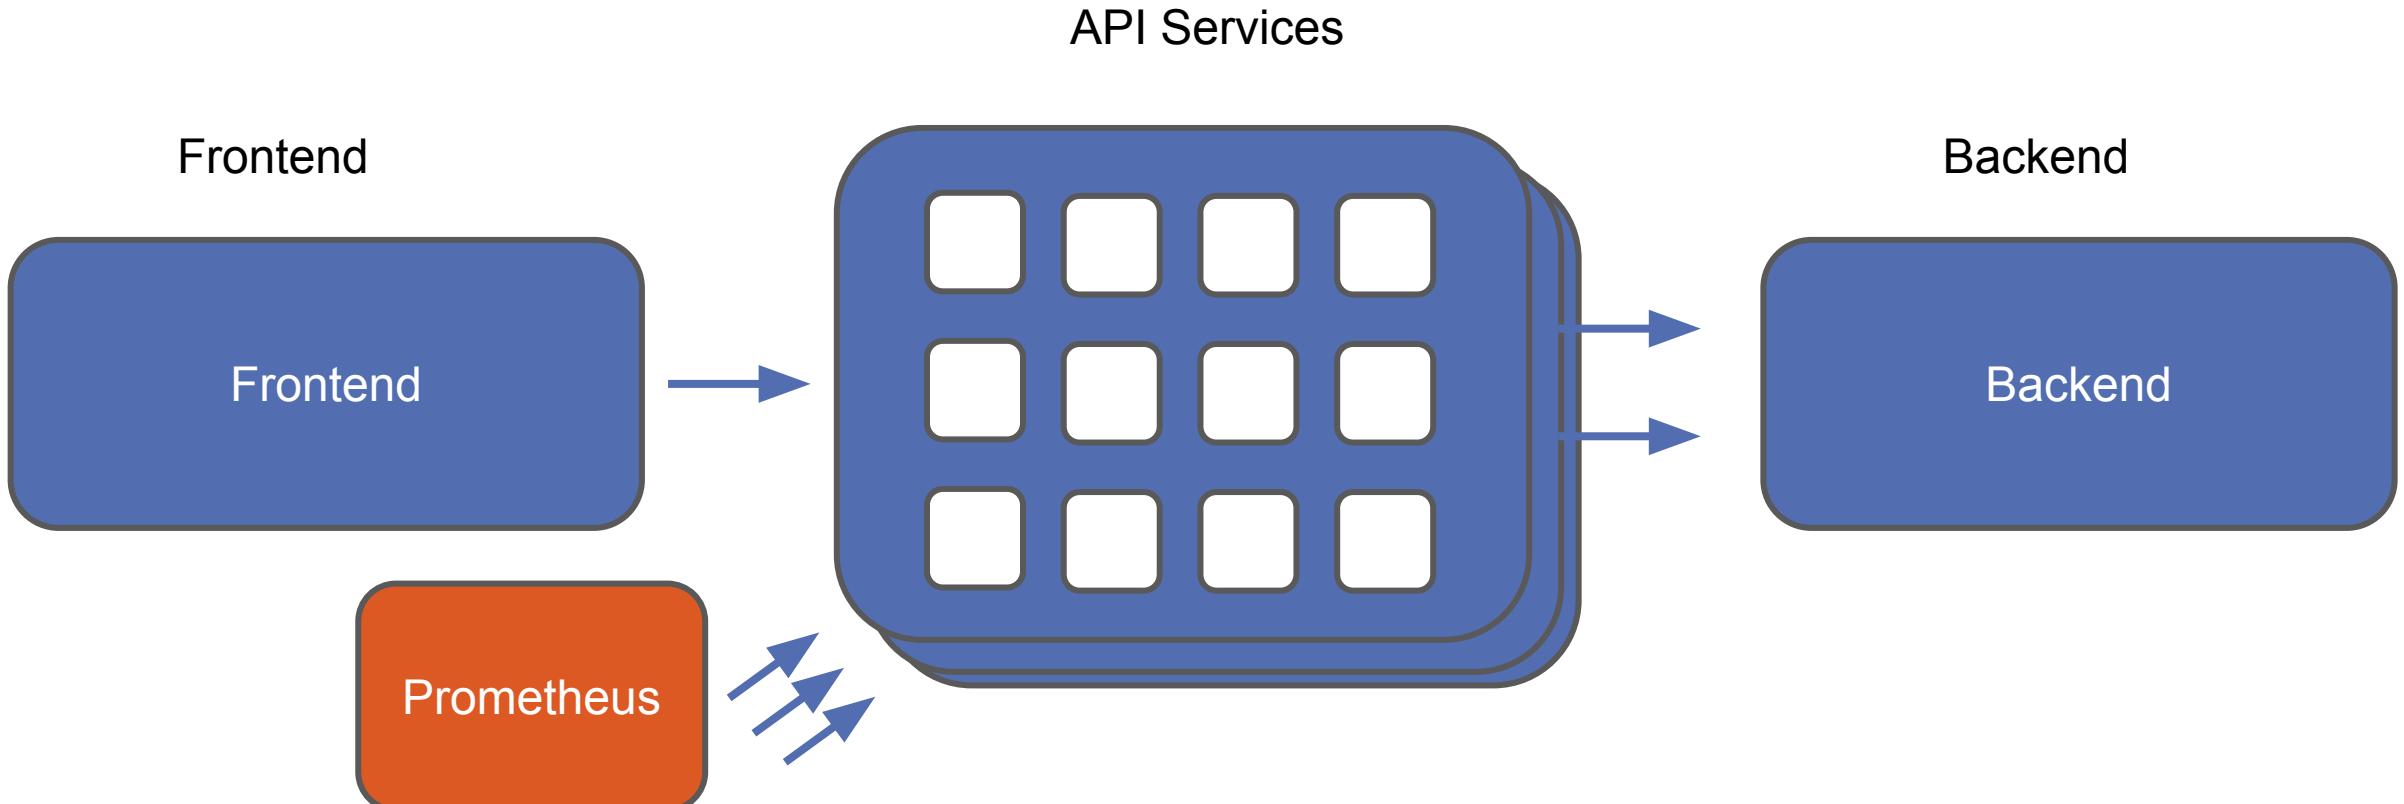
Limited  
visibility

# MONITORING



Monitor  
Internals

# MONITORING



Monitor  
Services

# MONITORING

# Prometheus

The screenshot shows the Prometheus web interface. At the top, there is a dark navigation bar with the following items: Prometheus, Alerts, Graph, Status ▾, and Help. Below the navigation bar is a search bar labeled "Expression (press Shift+Enter for newlines)". To the right of the search bar are two buttons: "Execute" (in blue) and "- insert metric at cursor -". Below the search bar is a tab bar with "Graph" selected (it is highlighted in blue) and "Console" as the other option. Underneath the tabs is a table with two columns: "Element" and "Value". The table contains a single row with the text "no data". At the bottom left of the interface is a blue button labeled "Add Graph".

# MONITORING

# Prometheus

Prometheus   Alerts   Graph   Status ▾   Help

## Targets

| kubernetes-apiservers (1/1 up)  |       |                              |             |       |
|---|-------|------------------------------|-------------|-------|
| Endpoint  | State | Labels                       | Last Scrape | Error |
| <a href="https://192.168.64.2:8443/metrics">https://192.168.64.2:8443/metrics</a> | UP    | instance="192.168.64.2:8443" | 49.262s ago |       |

| kubernetes-cadvisor (1/1 up)  |       |   |             |       |
|---|-------|---|-------------|-------|
| Endpoint  | State | Labels  | Last Scrape | Error |
| <a href="https://kubernetes.default.svc:443/api/v1/nodes/minikube/proxy/metrics/cadvisor">https://kubernetes.default.svc:443/api/v1/nodes/minikube/proxy/metrics/cadvisor</a> | UP    | beta_kubernetes_io_arch="amd64" beta_kubernetes_io_os="linux" instance="minikube" kubernetes_io_hostname="minikube" | 46.023s ago |       |

| kubernetes-nodes (1/1 up)   |       |   |             |       |
|---|-------|---|-------------|-------|
| Endpoint  | State | Labels  | Last Scrape | Error |
| <a href="https://kubernetes.default.svc:443/api/v1/nodes/minikube/proxy/metrics">https://kubernetes.default.svc:443/api/v1/nodes/minikube/proxy/metrics</a> | UP    | beta_kubernetes_io_arch="amd64" beta_kubernetes_io_os="linux" instance="minikube" kubernetes_io_hostname="minikube" | 8.864s ago  |       |

| kubernetes-pods (1/1 up)  |       |   |             |       |
|---|-------|---|-------------|-------|
| Endpoint  | State | Labels  | Last Scrape | Error |
| <a href="http://172.17.0.3:9090/metrics">http://172.17.0.3:9090/metrics</a> | UP    | instance="172.17.0.3:9090" kubernetes_namespace="monitoring" kubernetes_pod_name="prometheus-5947dcc755-97qjd" name="prometheus" pod_template_hash="1503877311" | 8.719s ago  |       |

# MONITORING

# Grafana

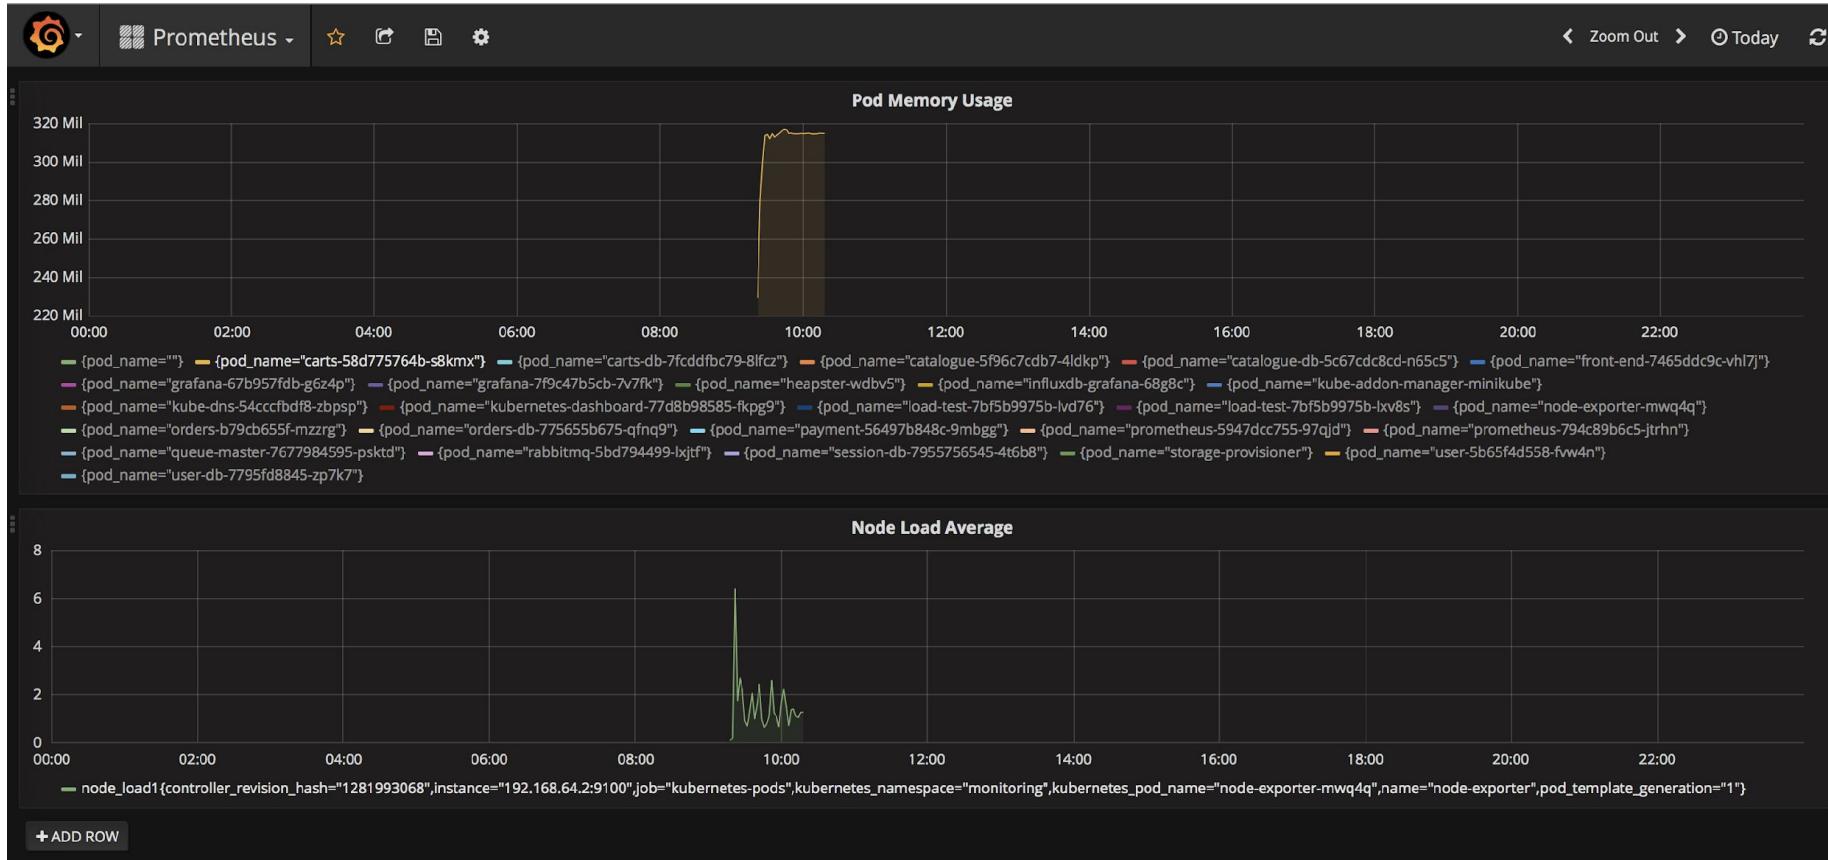
The screenshot shows the Grafana Home Dashboard. At the top, there is a navigation bar with icons for settings, home, and a gear. To the right are links for "Zoom Out", "Last 6 hours", and a refresh icon. The main title is "Home Dashboard". Below the title is a "Getting Started with Grafana" section featuring a five-step checklist:

- Step 1: "Install Grafana" (green checkmark)
- Step 2: "Create your first data-source" (green checkmark)
- Step 3: "Create your first dashboard" (green checkmark)
- Step 4: "Add Users" (green button)
- Step 5: "Install apps & plugins" (orange gear icon)

Below the checklist are two panels: "Starred dashboards" (listing "Prometheus") and "Recently viewed dashboards" (also listing "Prometheus"). To the right are sections for "Installed Apps" (none), "Installed Panels" (none), and "Installed Datasources" (none). Each section includes a "Browse Grafana.com" link.

# MONITORING

# Grafana



# MONITORING

# Grafana

The screenshot shows the Grafana Alerting interface. At the top, there is a navigation bar with a gear icon and the text "Alerting". Below the navigation bar, the title "Alert List" is displayed in a large, italicized font. To the right of the title are two buttons: "Configure notifications" and "How to add an alert". Below the title, there are two buttons: "Filter by state" and "All". The main area below these buttons is completely empty, indicating no alerts are currently listed.

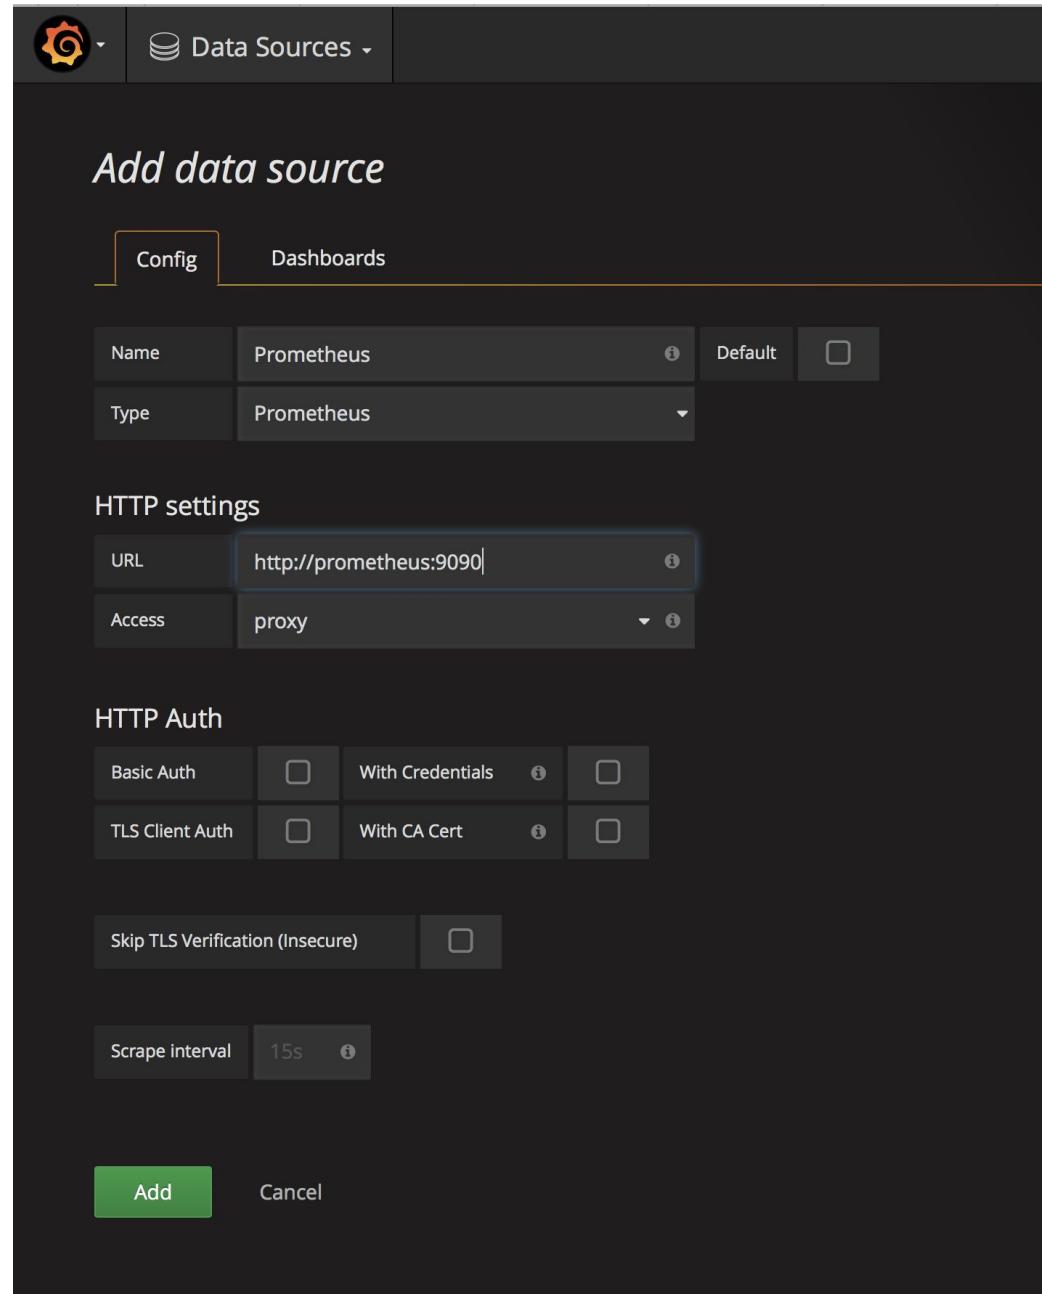
# MONITORING

# Grafana

The screenshot shows the Grafana interface for managing data sources. At the top, there is a navigation bar with a logo icon, a dropdown menu labeled "Data Sources", and a green button labeled "+ Add data source". Below the navigation bar, the title "Data Sources" is displayed in white. A horizontal line separates the title from the data source list. The first data source listed is "PROMETHEUS", which includes a Prometheus icon, the name "prometheus", a status indicator "default", and the URL "http://prometheus:9090". To the right of the data source list are two small icons: a grid icon and a list icon.

# MONITORING

# Grafana & Prometheus



# MONITORING

## Add graph

Graph      General      Metrics      Axes      Legend      Display      Alert      Time range      x

Data Source      prometheus      ▶ Query Inspector

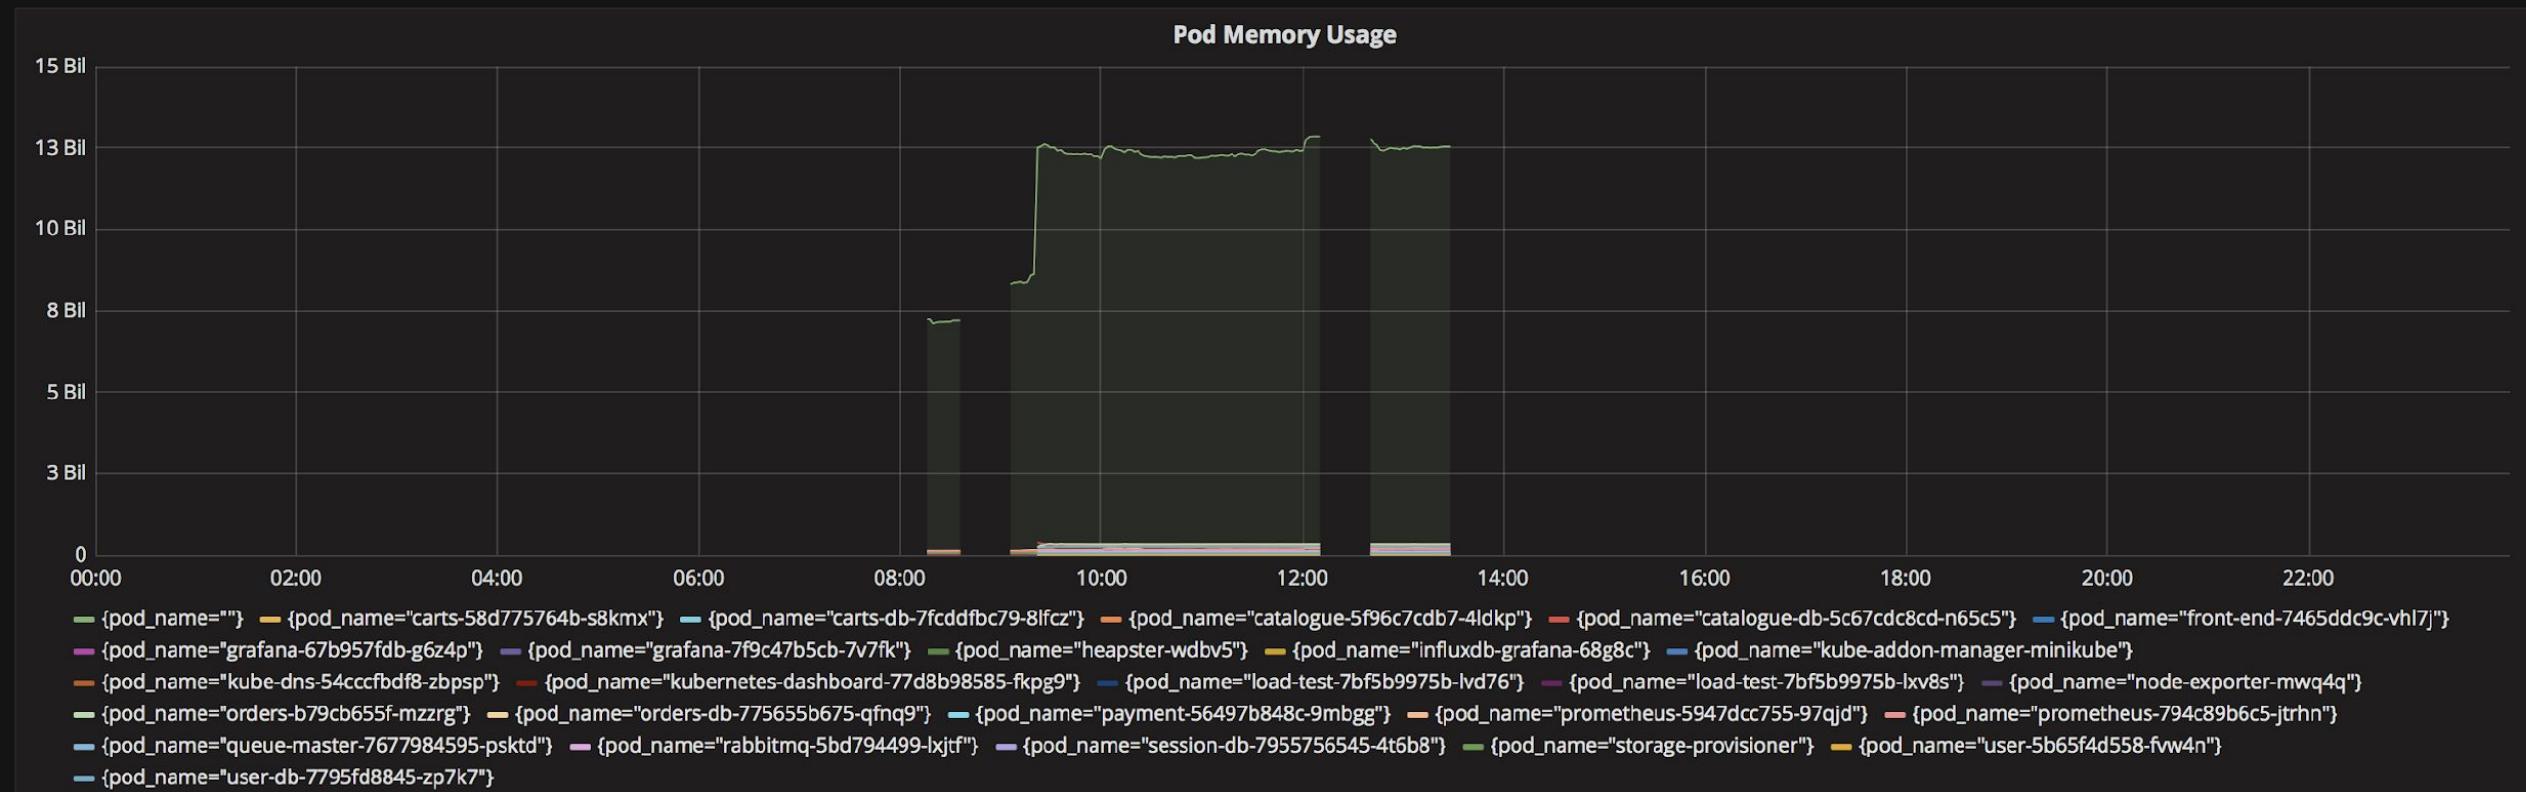
A      `sum(container_memory_usage_bytes) by (pod_name)`      ☰    ⚔    🗑

Legend format      legend format      Min step      1m      ⓘ      Resolution      1/2      Format as      Time series      Instant           

B      Add Query

# MONITORING

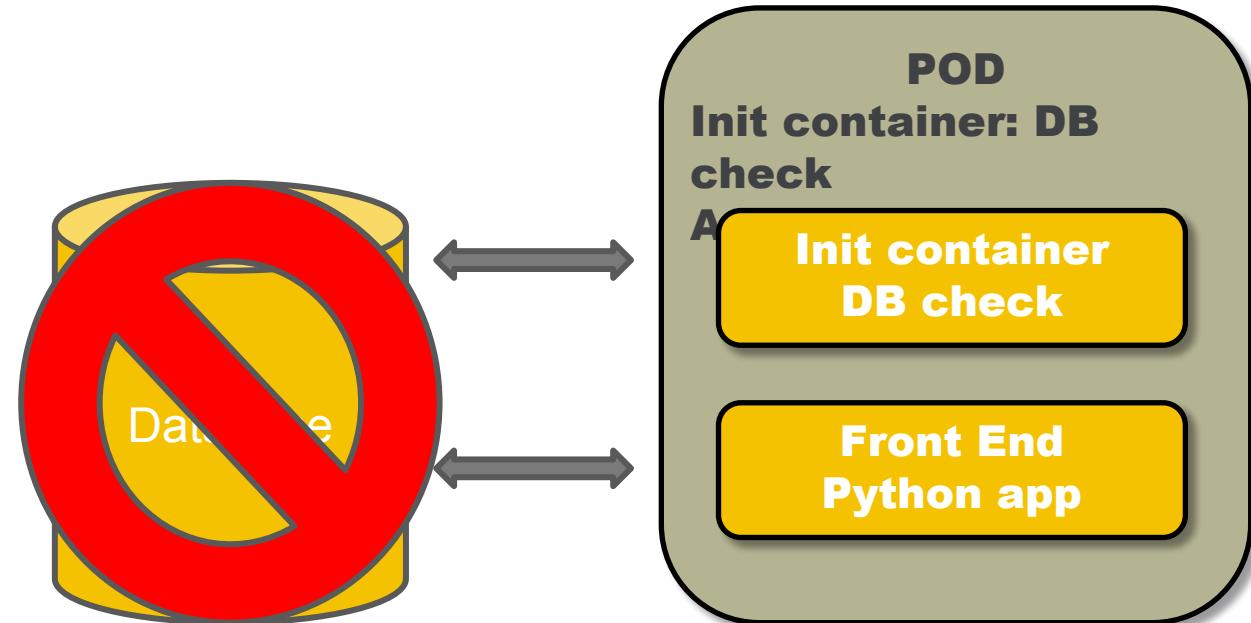
## Grafana graph



# INIT CONTAINERS

*Application orchestration*

- **Init container runs first to completion**
  - If failure, POD is restarted until successful.



# INIT CONTAINERS

**Support same features as application containers.**

- Resource limits
- Volumes
- Security settings
- Etc.

**Do not support readiness probes since they must run to completion before POD can be “ready”**

# KUBERNETES 1.5 SYNTAX

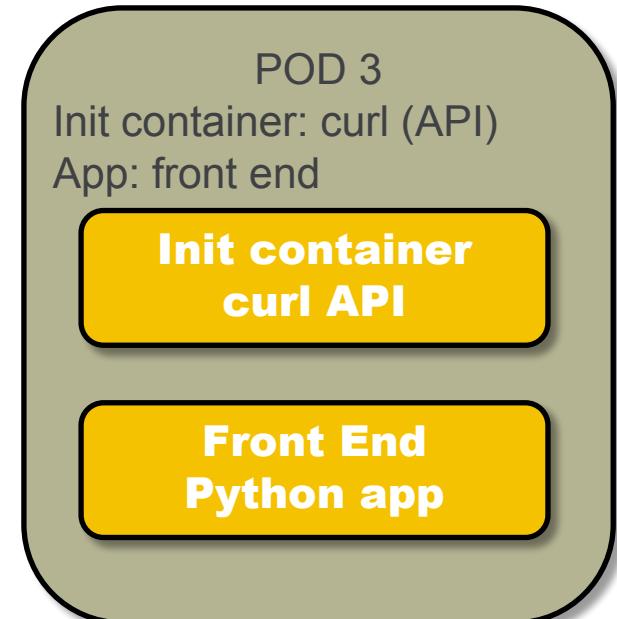
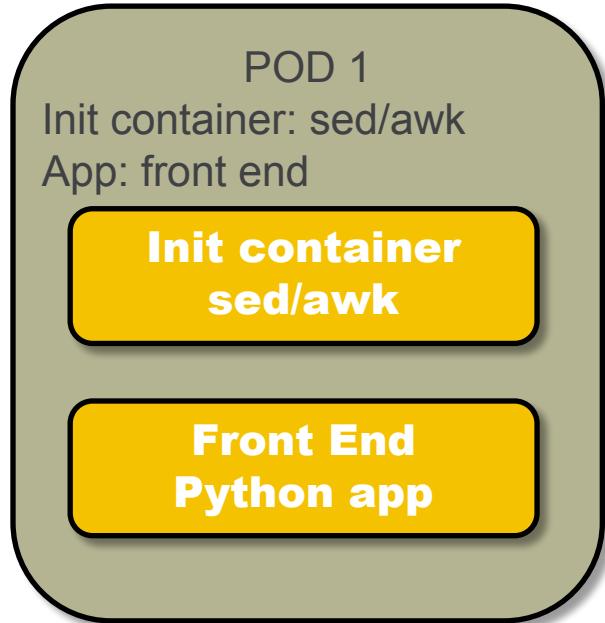
```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
  annotations:
    pod.beta.kubernetes.io/init-containers: '[{"name": "init-myservice", "image": "busybox", "command": ["sh", "-c", "until nslookup myservice; do echo waiting for myservice; sleep 2; done;"]}, {"name": "init-mydb", "image": "busybox", "command": ["sh", "-c", "until nslookup mydb; do echo waiting for mydb; sleep 2; done;"]}]'
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo The app is running! && sleep 3600']
```

# KUBERNETES 1.6+ SYNTAX

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo The app is running! && sleep 3600']
  initContainers:
  - name: init-myservice
    image: busybox
    command: ['sh', '-c', 'until nslookup myservice; do echo waiting for myservice; sleep 2; done;']
  - name: init-mydb
    image: busybox
    command: ['sh', '-c', 'until nslookup mydb; do echo waiting for mydb; sleep 2; done;']
```

# INIT CONTAINER USE-CASES

Lots of reasons



# STORING SECRETS INSIDE KUBERNETES



## What secrets do applications have?

- Database credentials
- API credentials & endpoints (Twitter, Facebook etc.)
- Infrastructure API credentials (Google, Azure, AWS)
- Private keys (TLS, SSH)
- Many more!

# STORING SECRETS INSIDE KUBERNETES

**It is a bad idea to include these secrets in your code.**

- Accidentally push up to GitHub with your code
- Push into your file storage and forget about
- Etc.

# STORING SECRETS INSIDE KUBERNETES

**There are bots crawling GitHub searching for secrets**

**Real life example:**

- Dev put keys out on GitHub, woke up next morning with a ton of emails and missed calls from Amazon
- 140 instances running under Dev's account.
- \$2,375 worth of Bitcoin mining.

# CREATE SECRET

- **Designed to hold all kinds of sensitive information**
- **Can be used by Pods (filesystem & environment variables) and the underlying kubelet when pulling images**

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: mmyWfoidfluL==
  username: NyhdOKwB
```

# POD SECRET

```
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: mycontainer
      image: redis
      env:
        - name: SECRET_USERNAME
          valueFrom:
            secretKeyRef:
              name: mysecret
              key: username
```

# VOLUME SECRET

```
spec:  
  containers  
    - name: mycontainer  
      image: redis  
      volumeMounts:  
        - name: "secrets"  
          mountPath: "/etc/my-secrets"  
          readOnly: true  
  volumes:  
    - name: "secrets"  
      secret:  
        secretName: "mysecret"
```



# **MICROSERVICES IN PRACTICE**

# **MICROSERVICES IN PRACTICE**

**Implementing a microservice architecture poses specific design, development and operational challenges.**

**Specifically, we want to address the following challenges:**

- Microservice Patterns
- Microservice Implementation Solutions



# MICROSERVICE PATTERNS

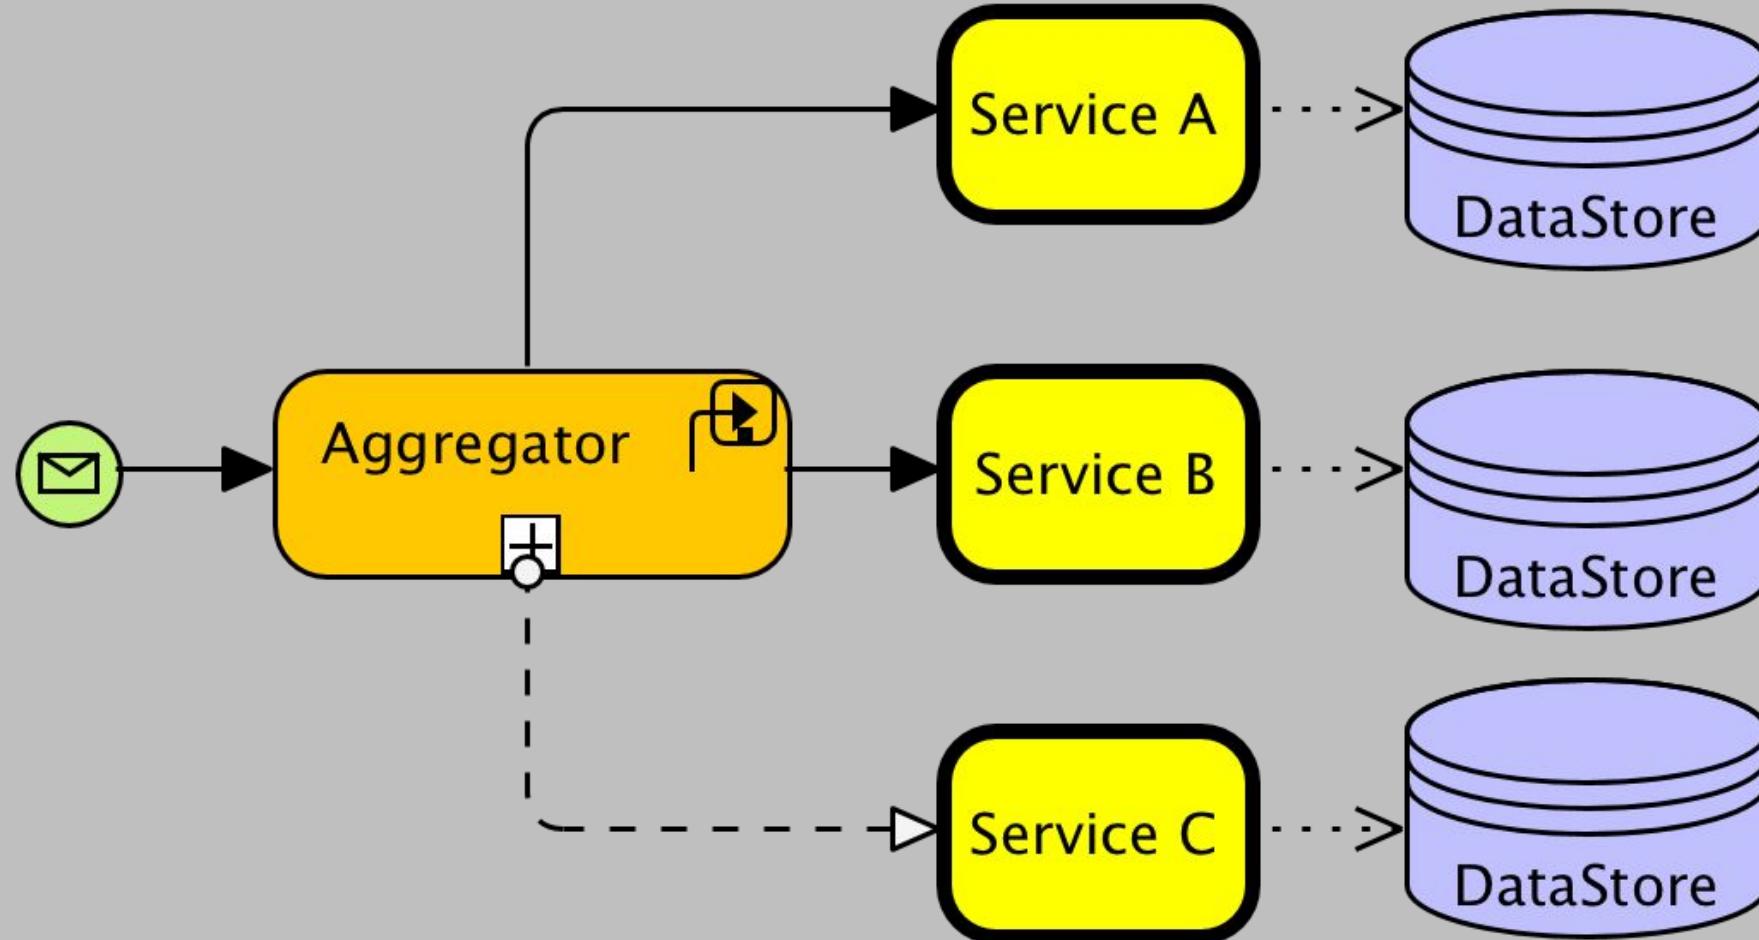
# MICROSERVICES PATTERNS

- Aggregator
- Branch
- Proxy
- Chained
- Shared Resources (Data)
- Asynchronous Messaging
- Bulkhead Isolation

# PATTERN: AGGREGATOR

## Description:

- An Aggregator is a web page that invokes multiple services to achieve the required functionality
- An Aggregator could also be a composite service
- This design pattern follows the DRY principle
- An Aggregator is similar to the Façade Pattern

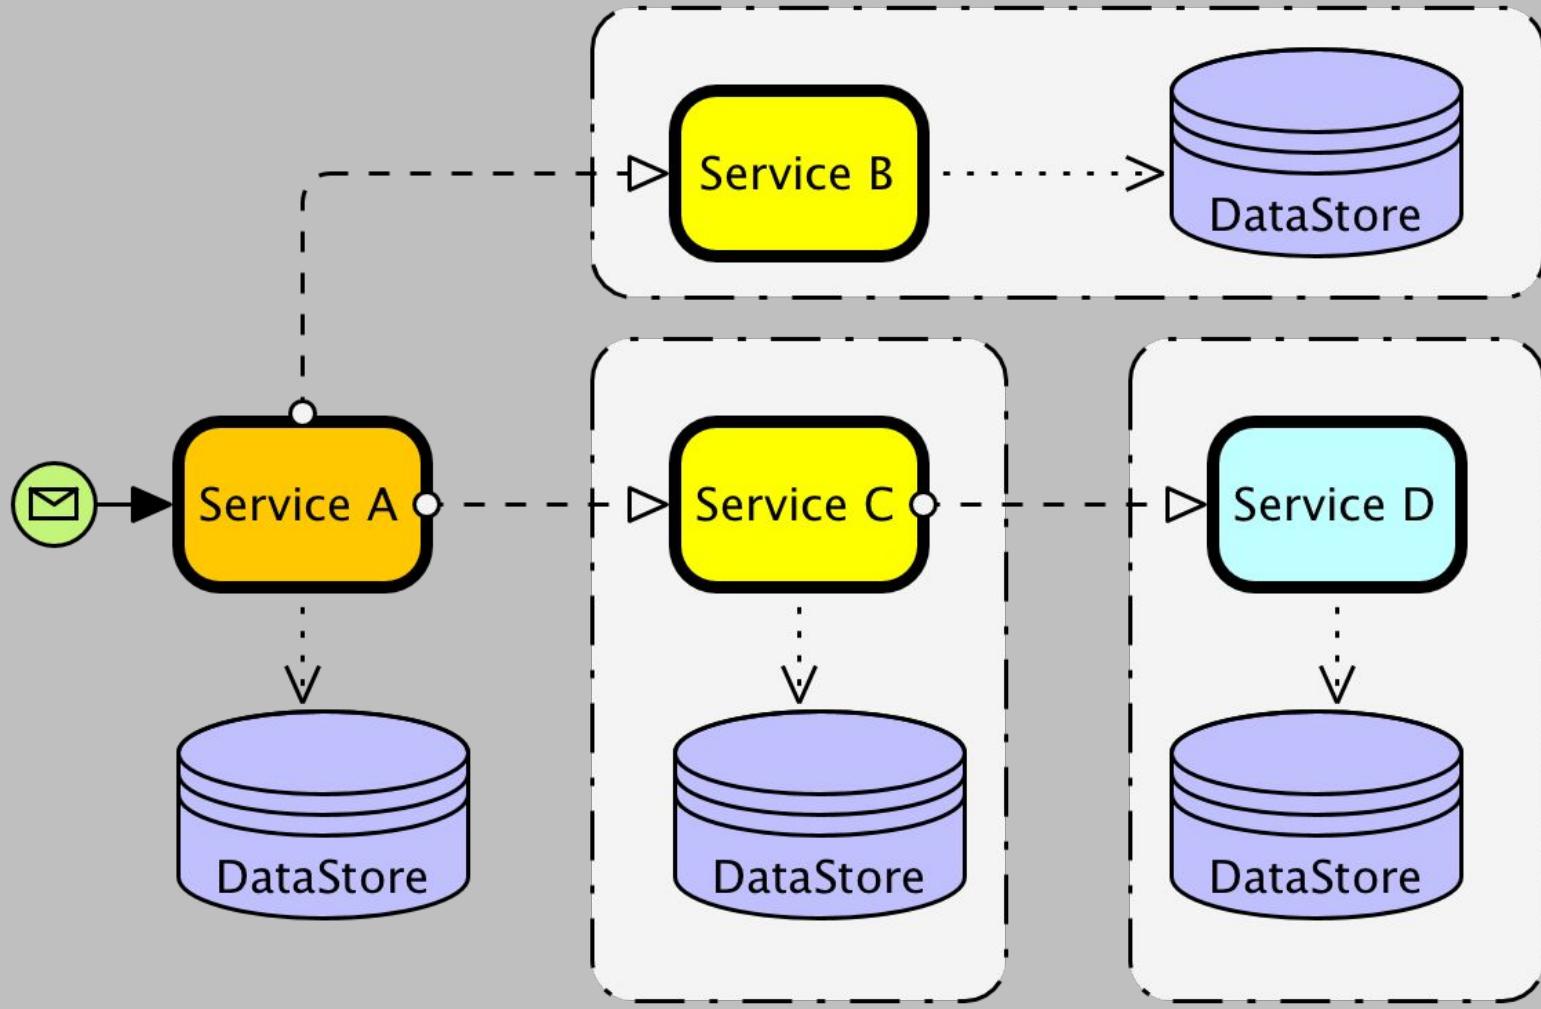


## PATTERN: AGGREGATOR

# PATTERN: BRANCH

## Description:

- The Branch design pattern extends the Aggregator design pattern
- Allows simultaneous response processing from two, likely mutually exclusive, chains of microservices
- This pattern can also be used to call different chains, or a single chain, based upon the business needs
- Service A, either a web page or a composite microservice, can invoke two different chains concurrently in which case this will resemble the Aggregator design pattern
- This may be configured using routing of EIP endpoints, and would need to be dynamically configurable

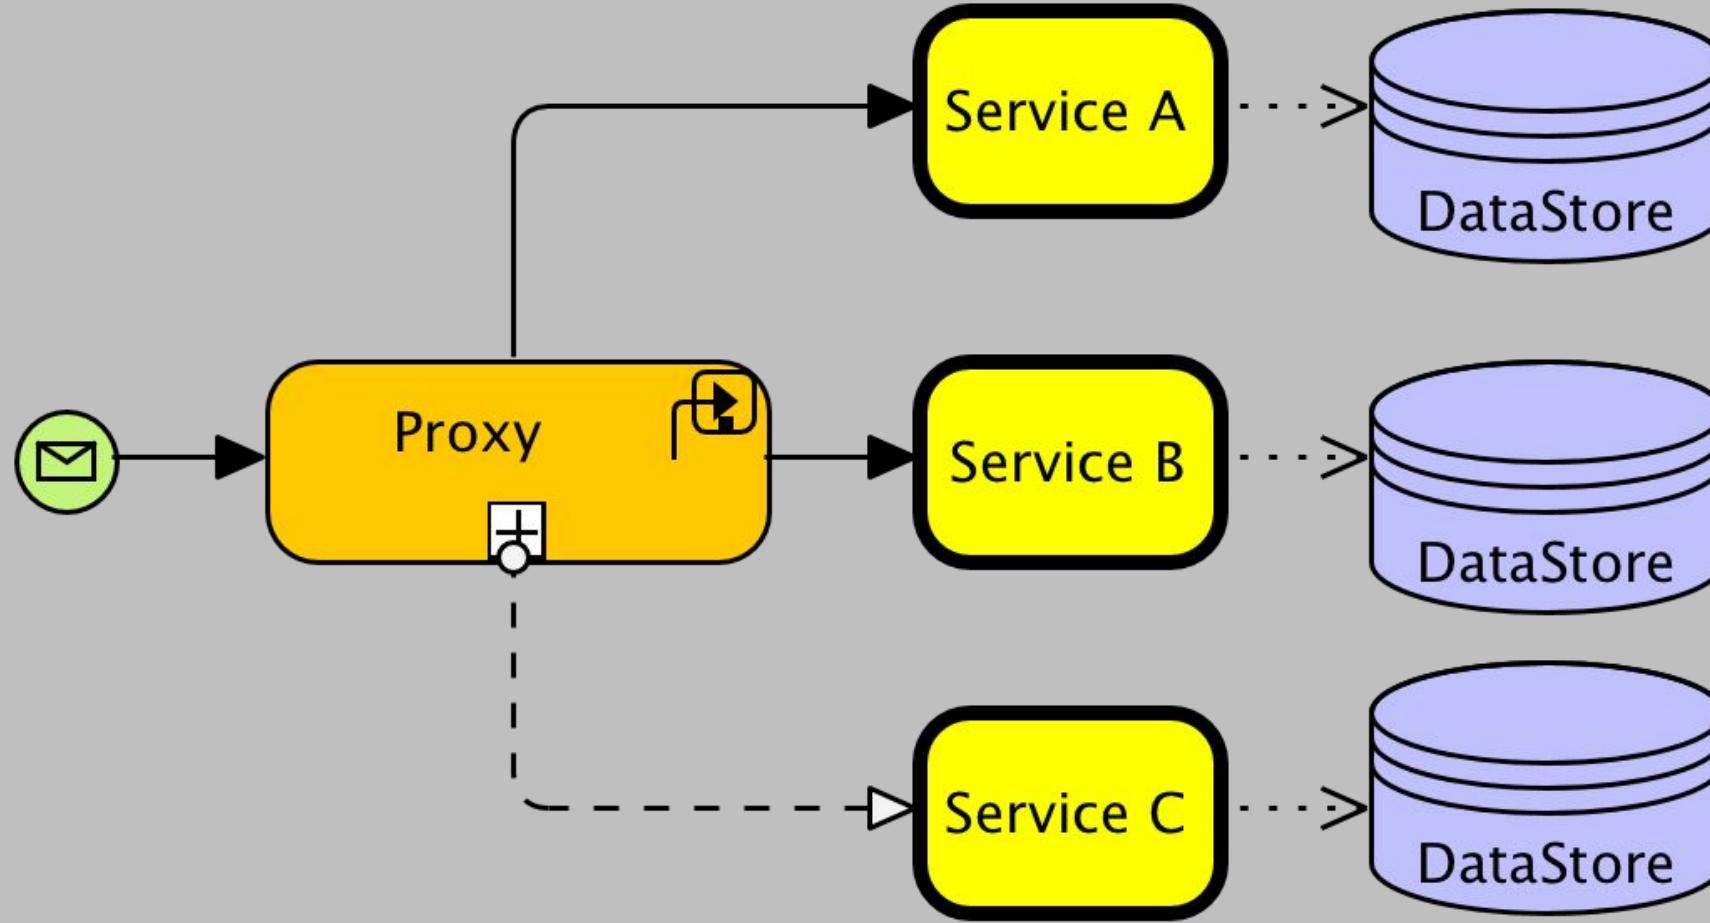


## PATTERN: BRANCHING

# PATTERN: PROXY

## Description:

- Proxy microservice design pattern is a variation of the Aggregator Pattern
- Different microservice may be invoked based upon the business needs
- Each Proxy can scale independently
- A proxy is a consumer facing interface in lieu of exposing individual services
- May be a dumb proxy that delegates requests to one of the services
- May be a Smart Proxy, where data is transformed before response returned to consumer

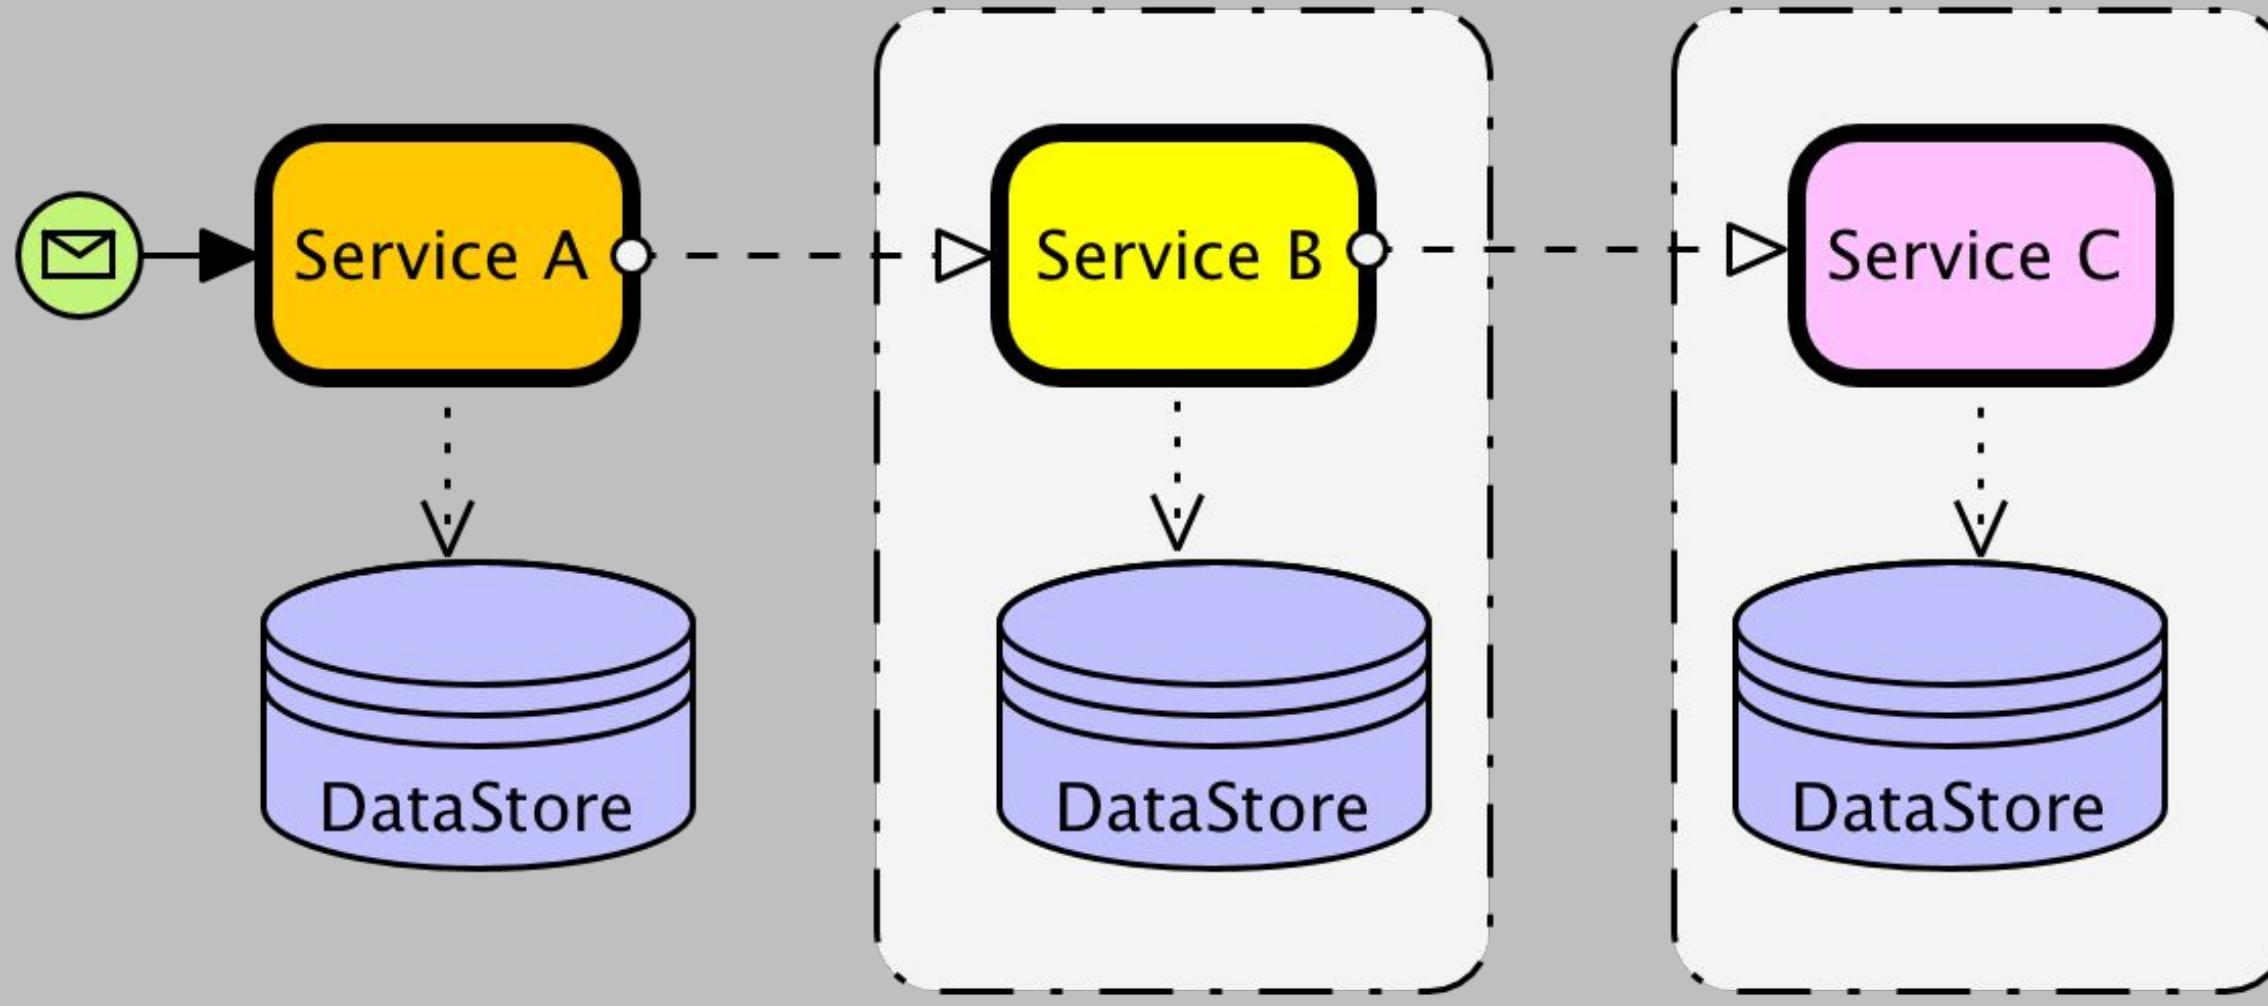


## PATTERN: PROXY

# PATTERN: CHAIN

## Description:

- Produces single consolidated response to request
- Request is received by Service A, which is chained to Service B, which in turn which is changed to Service C
- Services are likely using a synchronous HTTP request/response messaging
- Client is blocked until request/response chain is completed
- Important not to make chains too long
- A chain with a single microservice is called singleton chain

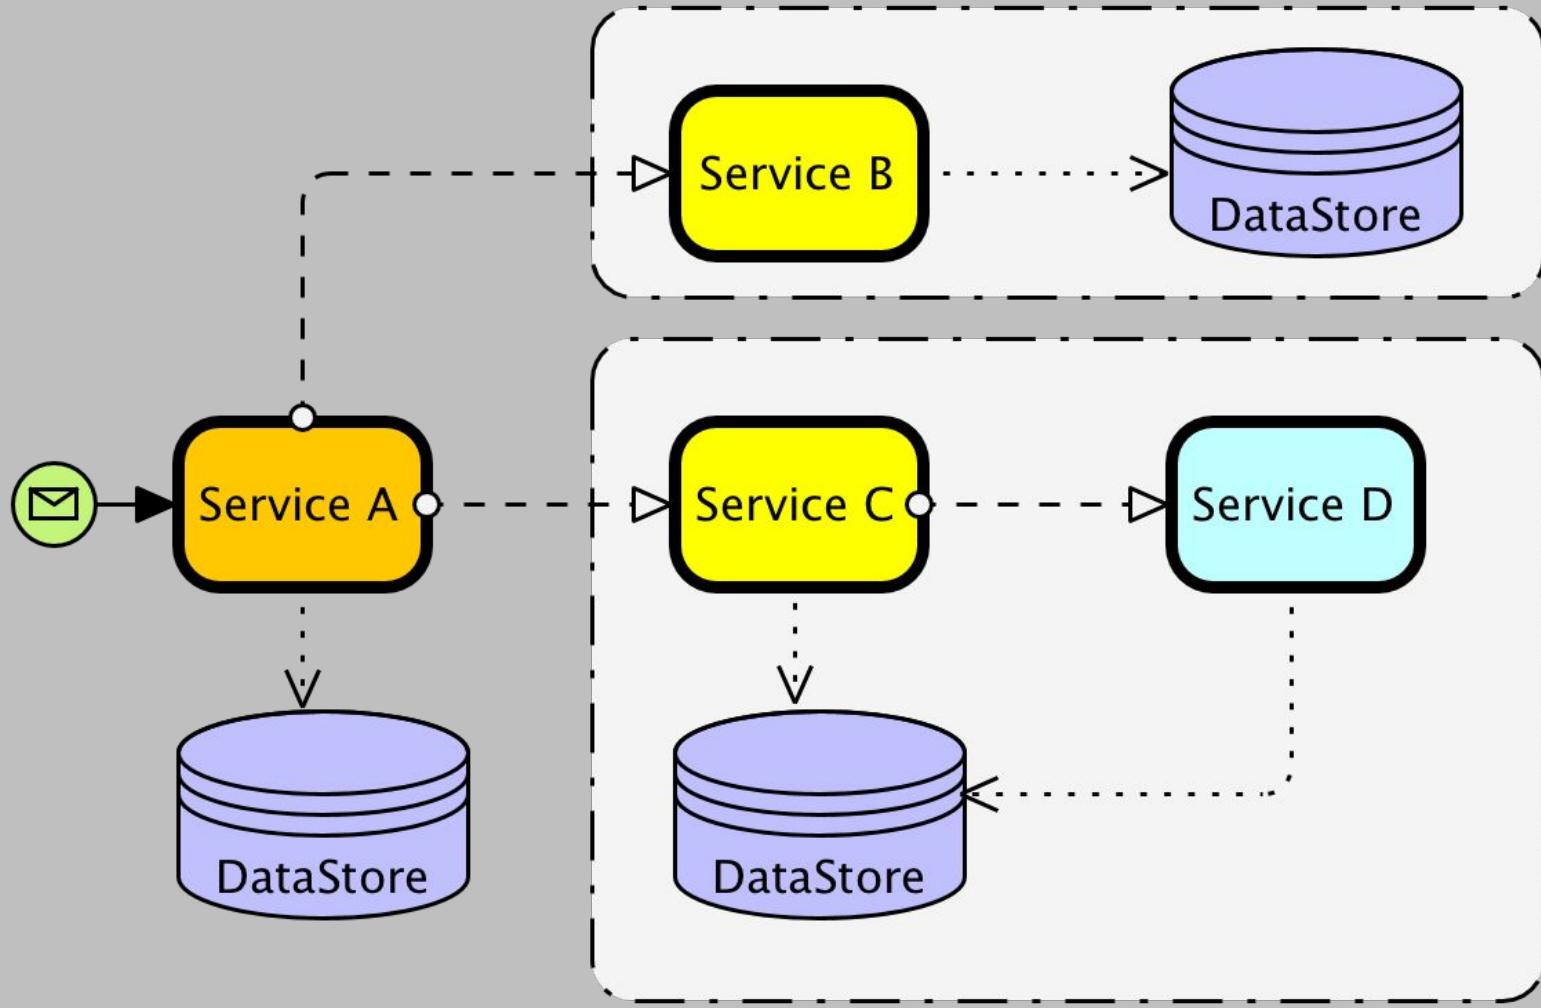


## PATTERN: CHAIN

# PATTERN: SHARED RESOURCES (DATA)

## Description:

- A fundamental design principle of microservice is autonomy
- A service should be full-stack and have control of all the components
  - UI, middleware, persistence, transaction, etc.
- Service should be polyglot
  - Use the right tool for the right job
- Data normalization is typical issue when refactoring existing monolithic applications
- This pattern may benefit from shared caching and database stores
- This is a transitional pattern until service can be fully autonomous

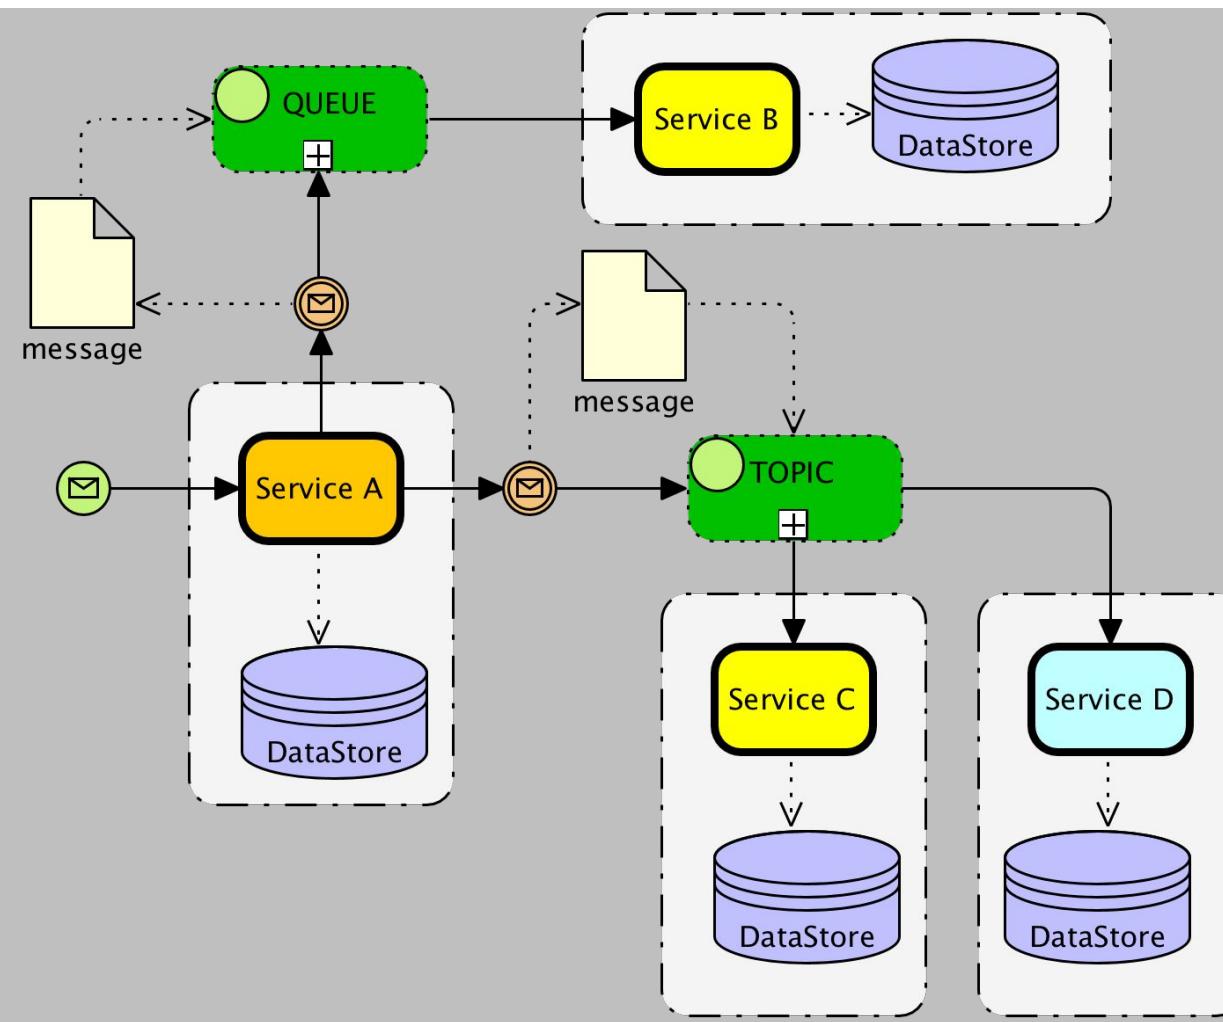


## PATTERN: SHARED RESOURCE (DATA)

# PATTERN: ASYNCHRONOUS MESSAGING

## Description:

- RESTful services are inherently synchronous, and thus blocking
- Microservice architectures may use asynchronous messaging instead of RESTful request/response
- A service may call another service synchronously, which then communicates with other services asynchronously
- Could communicate asynchronously in many different forms:
  - WebSockets, Message Queues, Reactive Streams
- Can combine RESTful service and publish/subscribe messaging to accomplish the business needs

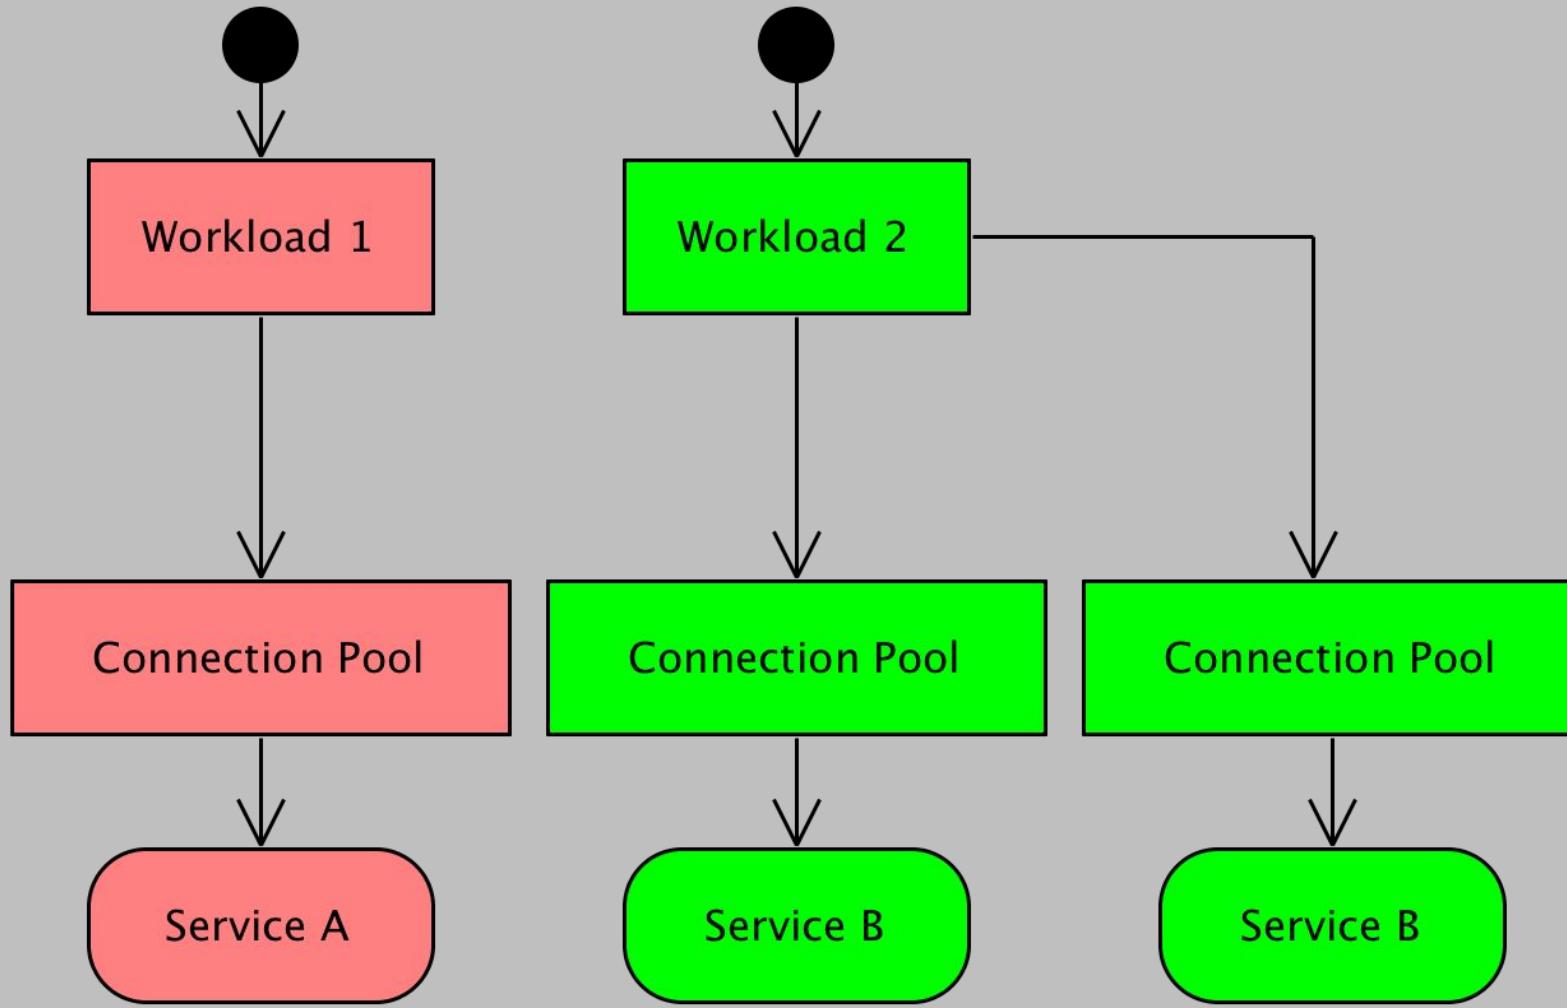


## PATTERN: ASYNCHRONOUS MESSAGING

# PATTERN: BULKHEAD ISOLATION

## Description:

- Isolate elements of an application into pools so that if one fails, the others will continue to function
- This pattern resembles the sectioned partitions of a ship's hull
  - If ship hull is compromised, only the damaged section fills with water, which prevents the ship from sinking
- Partition service instances into different groups, based on consumer load and availability requirements
- Isolating failures, and allowing application to sustain service functionality for some consumers, even during a failure



## PATTERN: BULKHEAD

# PATTERN: BULKHEAD ISOLATION

## Issues and considerations

- Define partitions around the business and technical requirements of the application
- Consider level of isolation offered by the technology as well as the overhead in terms of cost, performance and manageability
- Consider combining bulkheads with retry, circuit breaker, and throttling patterns to provide more sophisticated fault handling
- Use processes, thread pools, and semaphores
- Use frameworks for creating consumer bulkheads
- Deploy into separate VM's, containers, or processes
- Isolate asynchronous service into different queues
- Monitor each partition's performance and SLA



# DATA ISLAND'S

# DATA ISLAND'S

- Microservices are not isolated islands, they come in systems
- Data systems are about exposing data. Services are about hiding data
- The more shared data hidden inside a service boundary, the more complex the interface
- The more complex the interface the harder it will be to join datasets across different services
- Extracting whole dataset into mutable islands, the more they will diverge over time
- Better to take a decentralized approach to data and services with data-streams such as Kafka

# **IMMUTABILITY**

## **Definition**

**Unchanging over time or unable to be changed, fixed, permanent, carved in stone**

# IMMUTABILITY

Once you created an artifact, be that a container image, or a VM image, or maybe a package from compiled code - you declare that you will never ever change it.

Often if any changes are required, you declare that a new version of "thing" will be created instead.

# **IDEMPOTENCE**

## **Definition**

**The result of a successful performed request is independent of the number of times it is executed.**

# IDEMPOTENCE

## NATURAL IDEMPOTENCY

Some services can be executed as often as you want because they just alter state.

Example:

*confirmCustomer()*

## BUSINESS IDEMPOTENCY

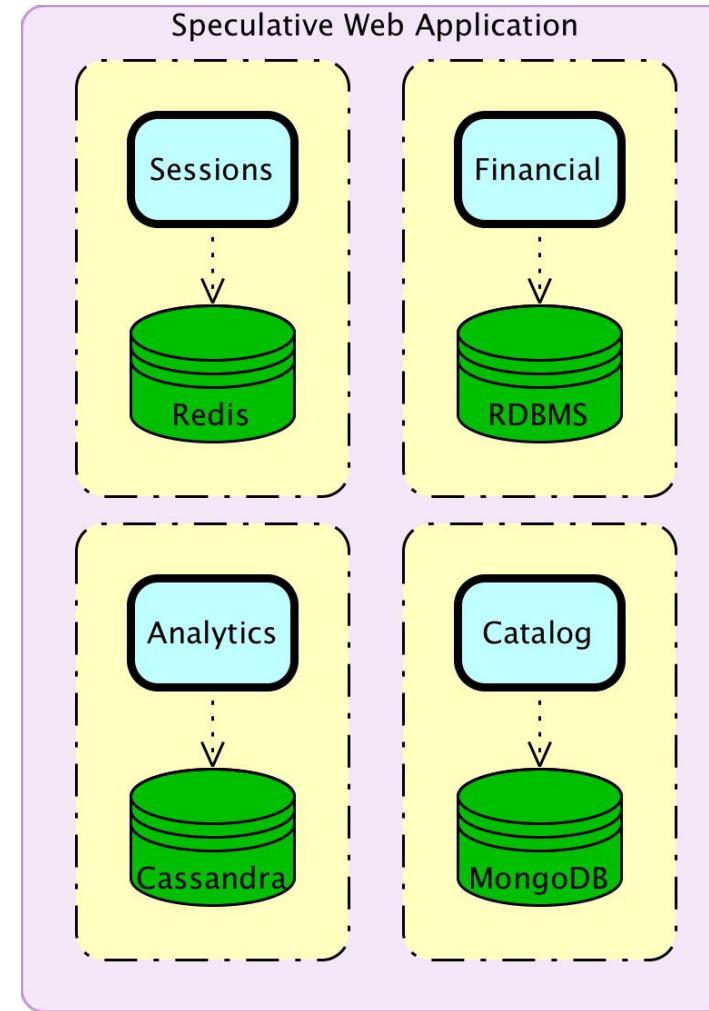
Sometimes business identifiers allow services to detect duplicate requests.

Example:

*createCustomer(email)*

# POLYGLOT PERSISTENCE

A significant advantage of decentralized data management is the ability to take advantage of polyglot persistence



# POLYGLOT PERSISTENCE

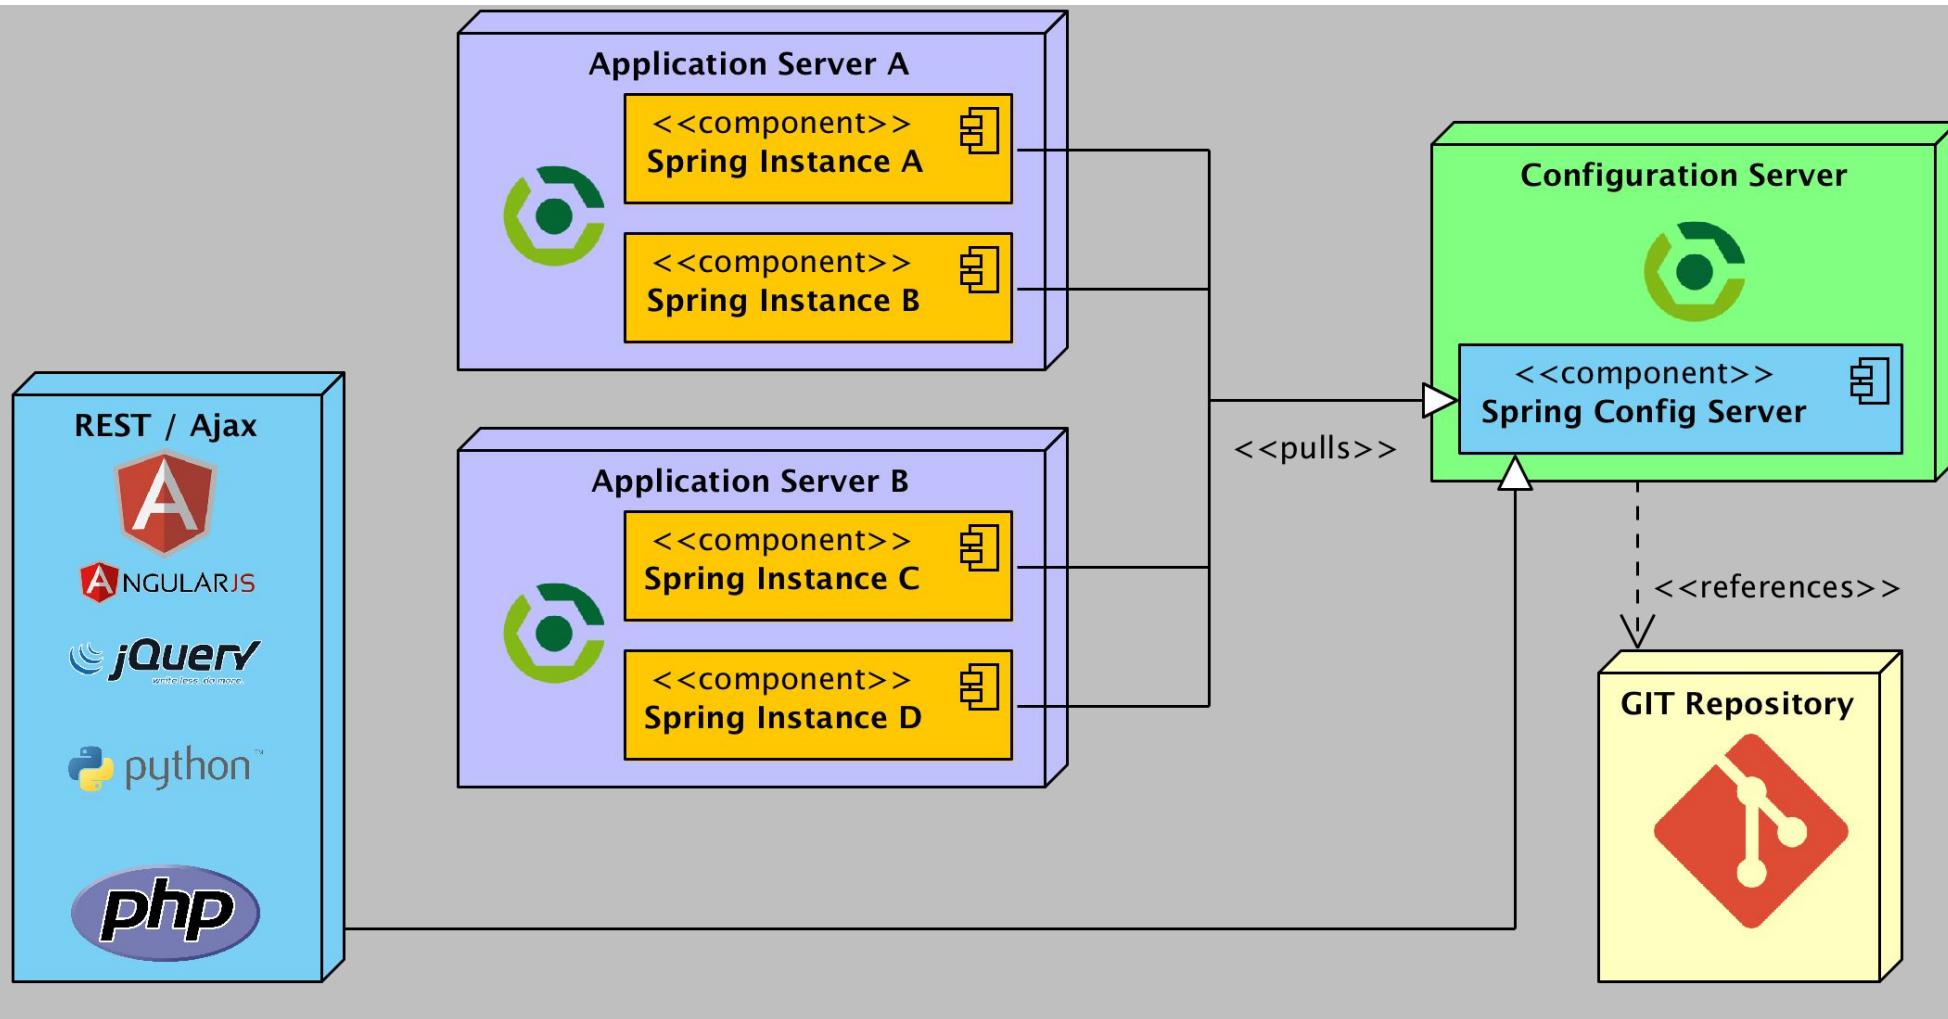
## Different types of data have different storage requirements:

- Read/Write Balance (Some types of data have a very high write volume. This can require a different type of data store compared to data that has low write volume but high read volume.)
- Data Structure (Some types of highly structured data such as JSON documents may be better stored in a NoSQL database such as MongoDB, while flat relational objects may be more efficiently stored in a SQL database.)
- Data Querying (Some data may be accessible using a simple key value store, while other types of data may require advanced querying based on the values of multiple columns.)
- Data Lifecycle (Some data is temporary in nature and can be stored in a fast, in-memory store such as Redis or Memcached, while other data must be retained for all time and needs very durable storage on disk.)
- Data Size (Some data is made up of fairly uniform rows of consistent byte size, while other data may include large blobs that need to be stored in something like AWS S3.)

# MICROSERVICE CONFIGURATION

**Microservices require externalized configuration as per the 12-factor app requirement to inject different configuration for different environments.**

- Configuration service is another service in the mix
- Configuration clients can request configuration based on
  - Service Name
  - Service Type
  - Service Environment
  - Service Version



# MICROSERVICE CONFIGURATION

# API GATEWAYS

**A single point of entry into the system.**

**The API is used to encapsulate the entire system and provide only the required API for the service to the client.**

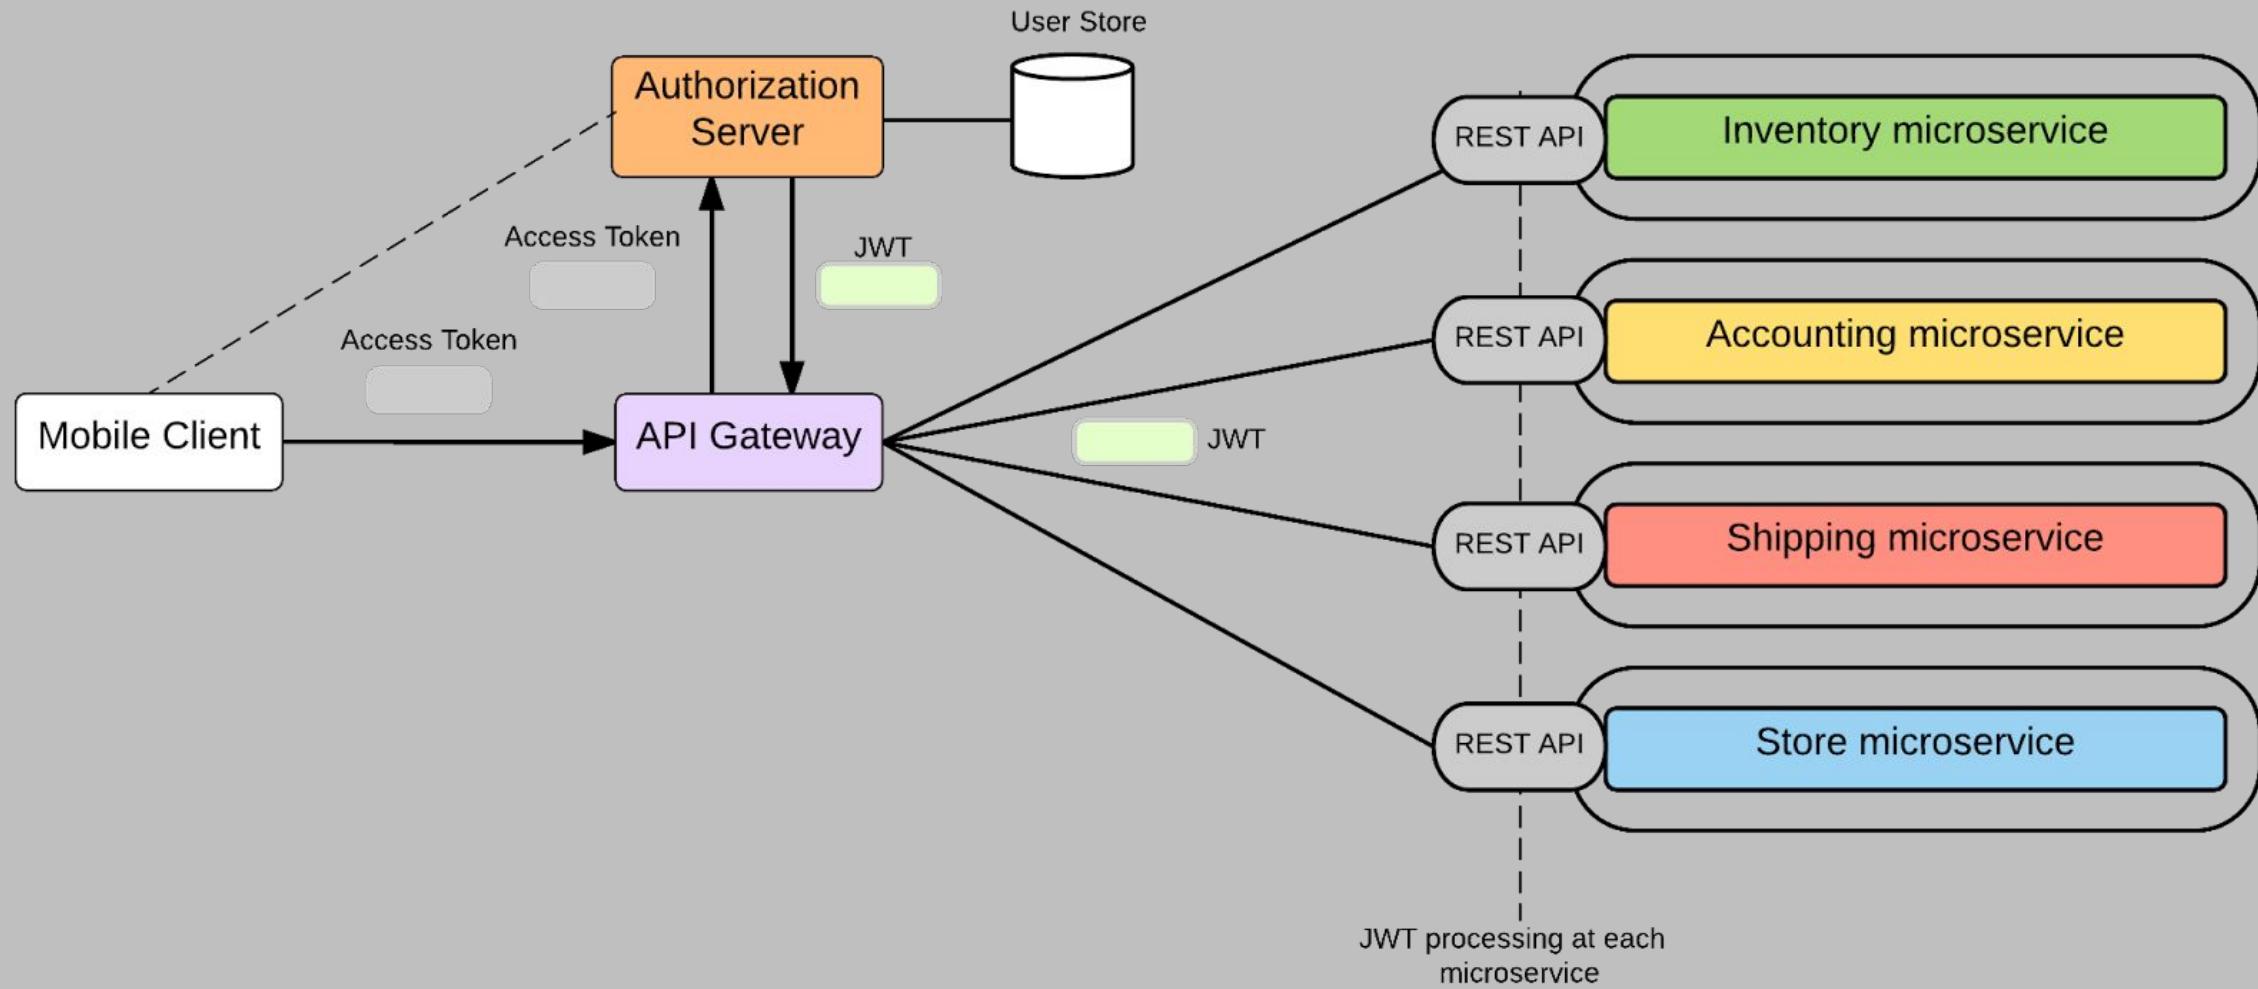
**The nature of the API definition depends on which IPC mechanism you are using.**

- If you are using messaging, the API consists of the message channels and the message types.
- If you are using HTTP, the API consists of the URLs and the request and response formats.

# API GATEWAYS (CONT'D)

**It might have other responsibilities such as:**

- Authentication
- Monitoring
- Load balancing
- Caching
- Static response handling
- More...



# API GATEWAY

# API GATEWAY CONSIDERATIONS

## Benefits

- A major benefit of using an API Gateway is that it encapsulates the internal structure of the application.

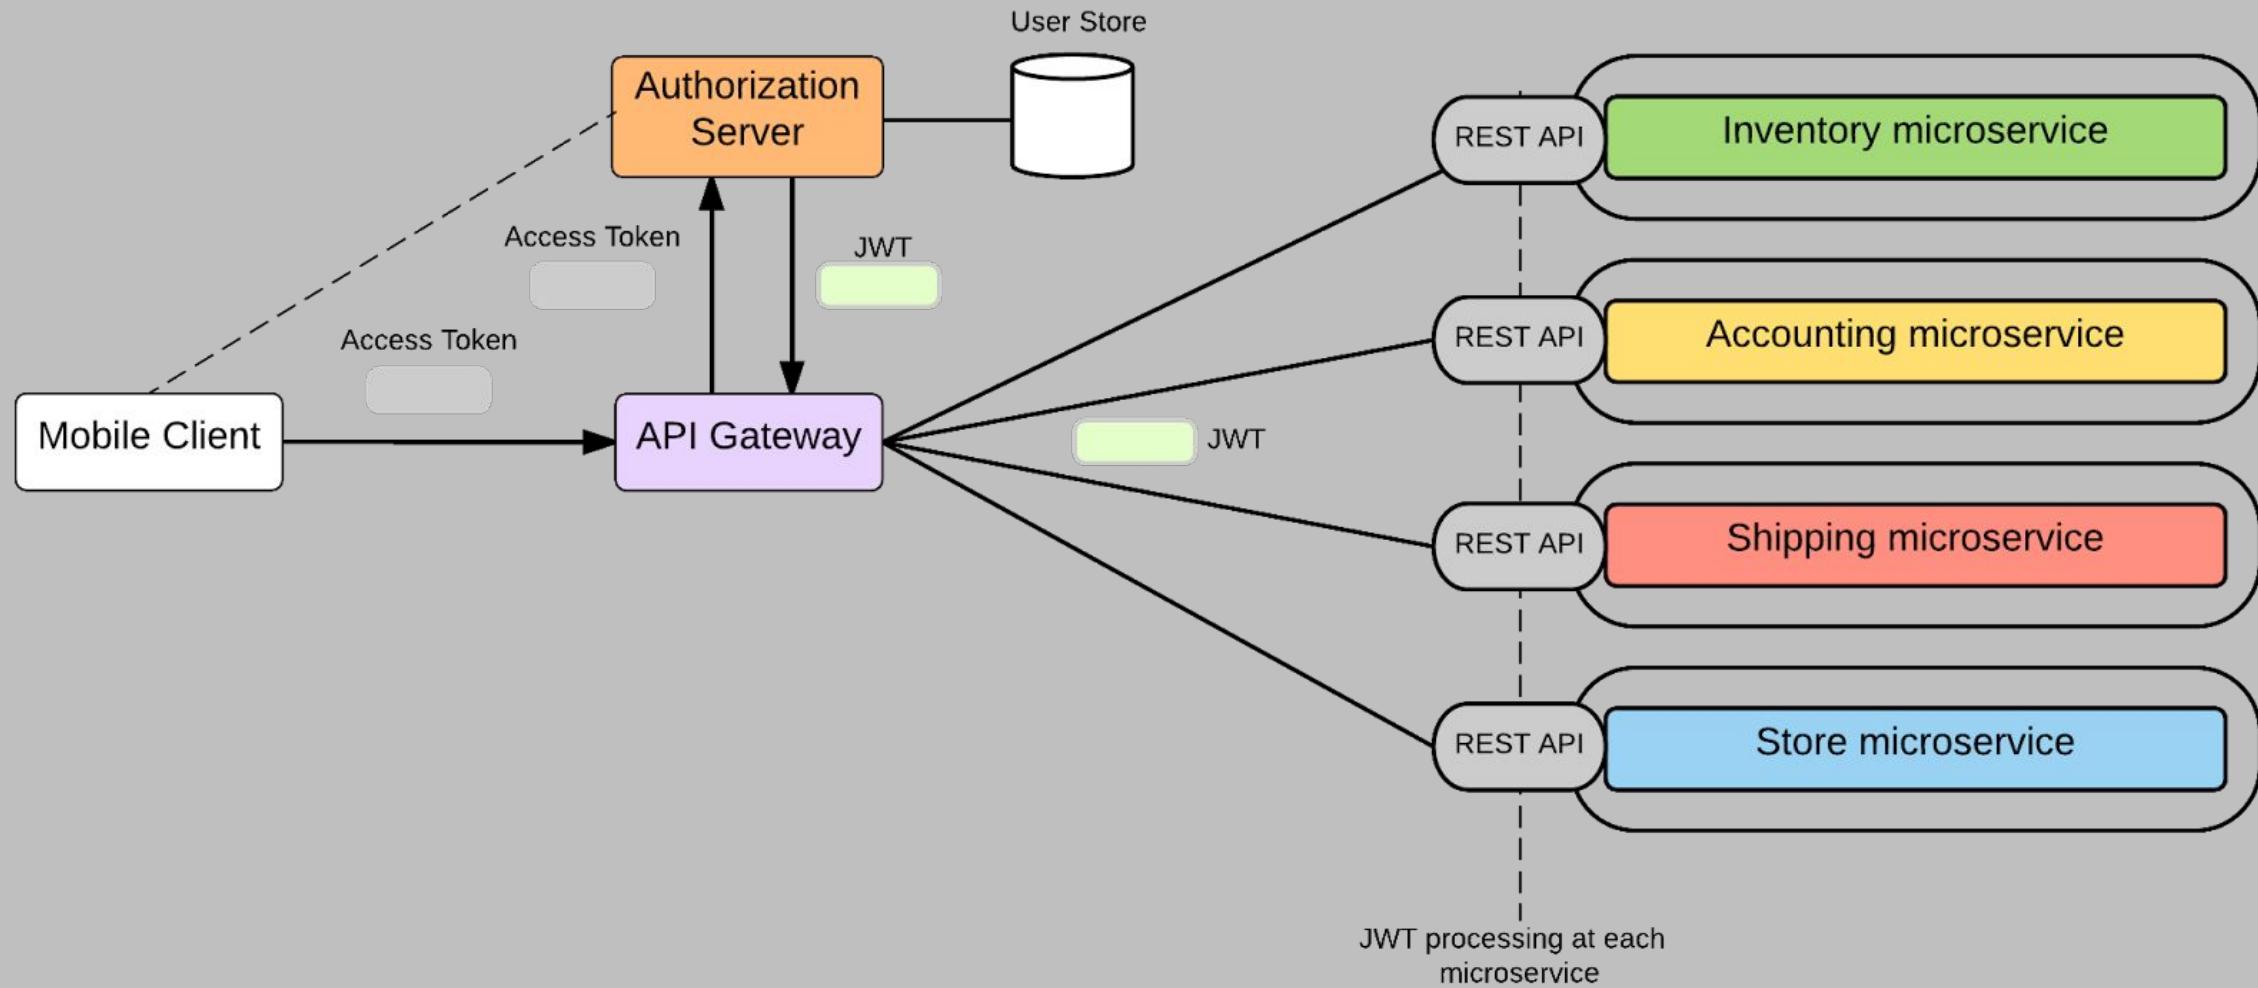
## Drawbacks

- There is also a risk that the API Gateway becomes a development bottleneck.
- Developers must update the API Gateway in order to expose each microservice's endpoints.

# **API GATEWAY (AUTHENTICATION)**

## **Authentication**

**API Gateways usually handle the authentication and authorization from the external callers to the microservice level.**



# API GATEWAY

# **API GATEWAY (AUTHENTICATION)**

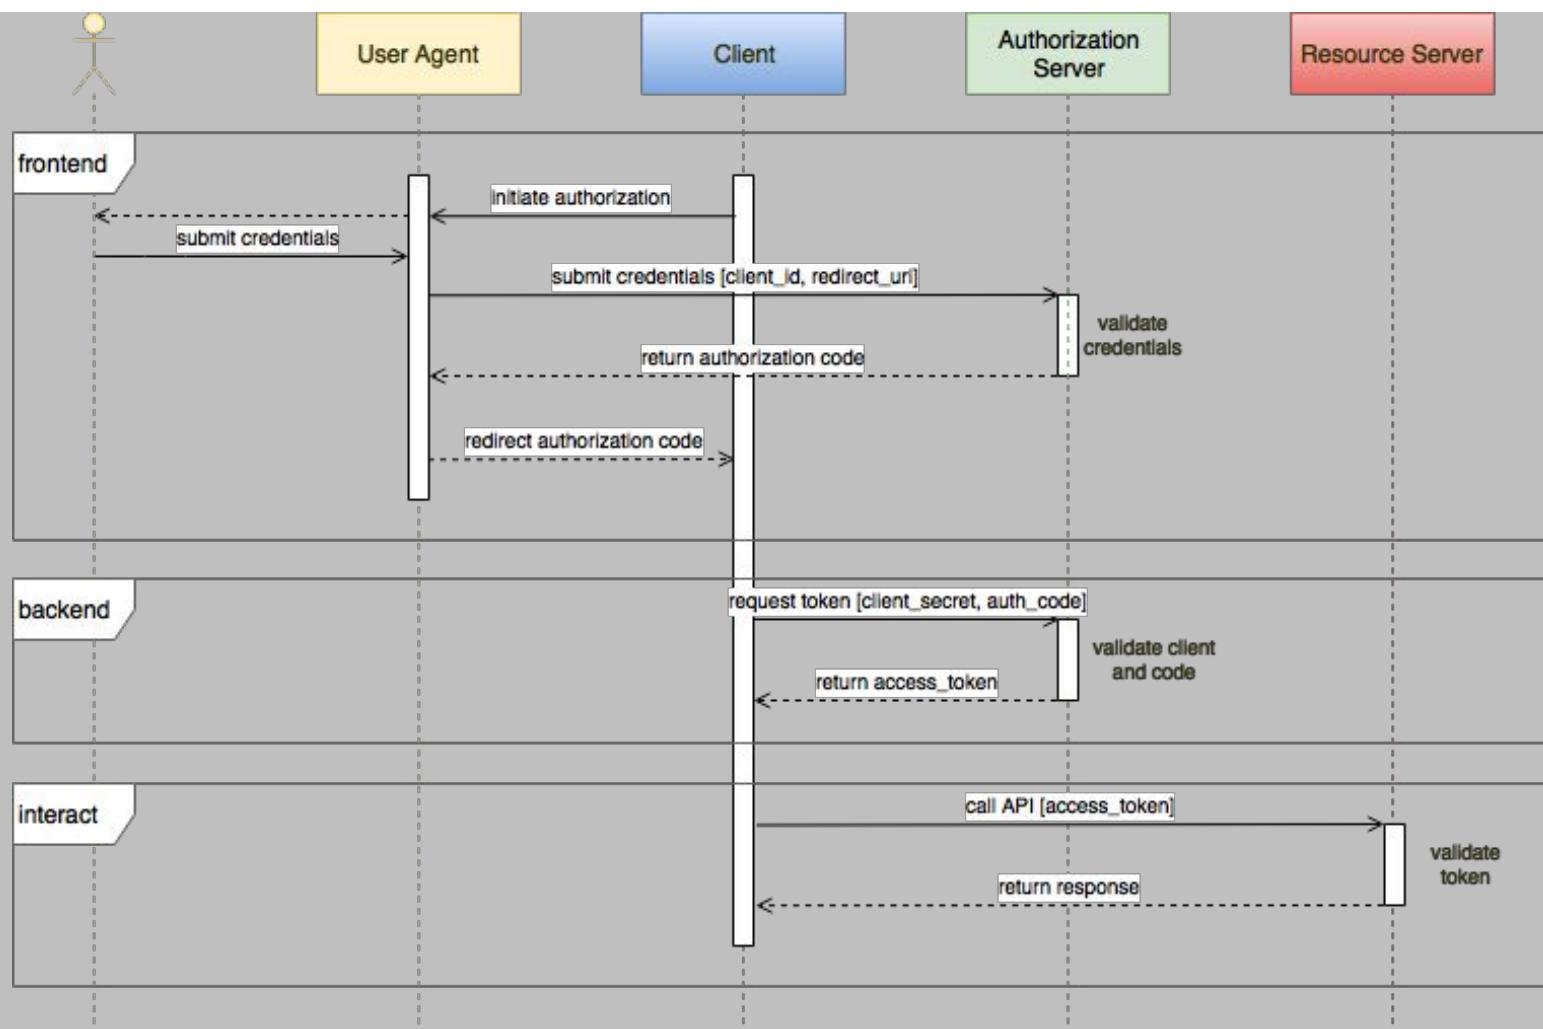
OAuth 2.0

**OAuth delegated authorization with JSON Web Tokens (JWT) is an efficient and scalable solution for authentication and authorization for microservices.**

# **API GATEWAY (AUTHENTICATION)**

## **JSON Web Token (JWT)**

**A JSON web token, or JWT (“jot”) for short, is a standardized, optionally validated and/or encrypted container format that is used to securely transfer information between two parties.**



# API GATEWAY (OAUTH)

OAuth and JWT Authentication

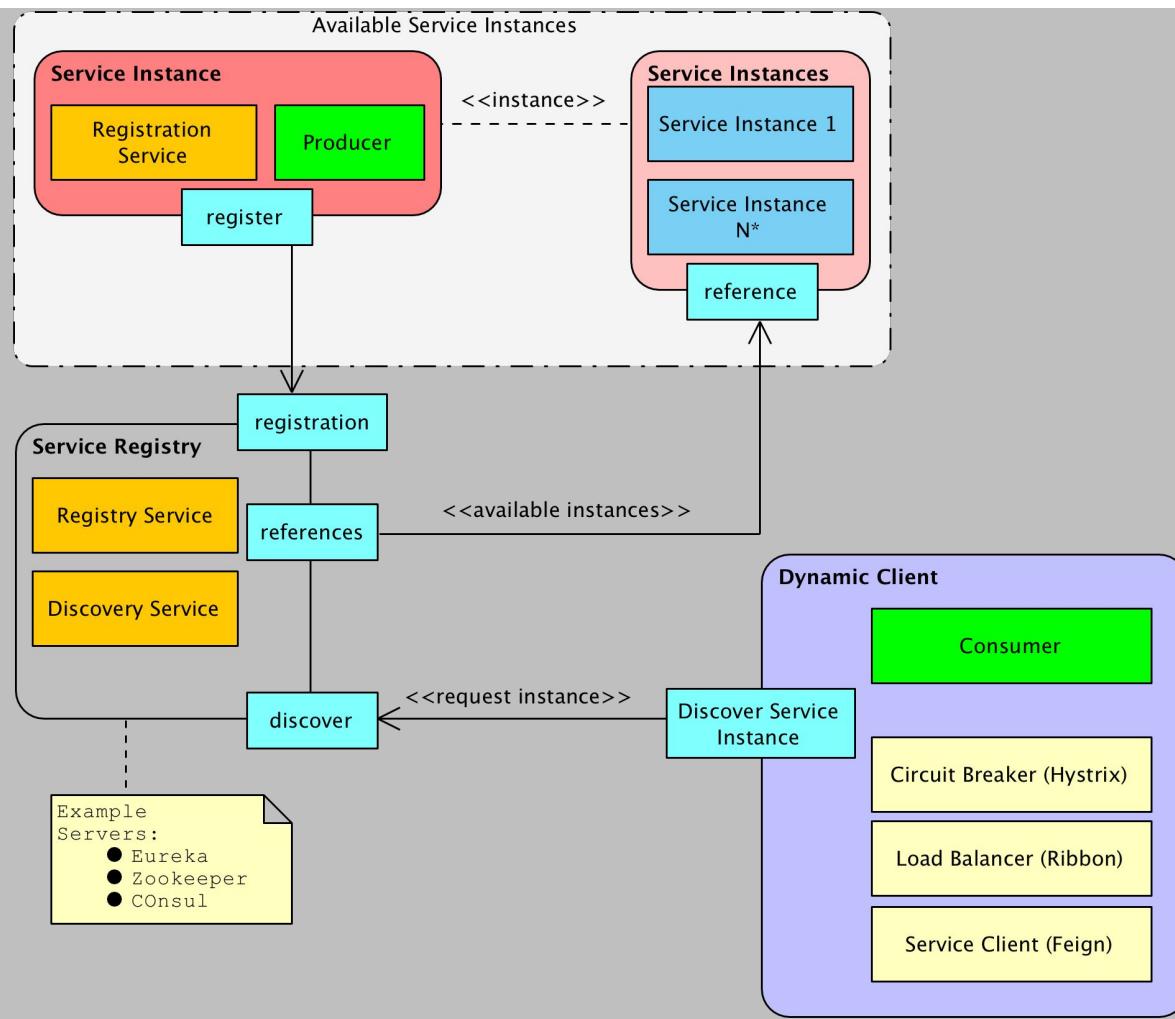
# API PROXYING

In simple terms you can think of the Proxy providing you a new endpoint for an existing API

Adds some lightweight security, monitoring, etc.

# MICROSERVICE REGISTRY / DISCOVERY





# MICROSERVICE REGISTRY / DISCOVERY

# MICROSERVICE REGISTRY / DISCOVERY

## Context

Clients of a service use either Client-side discovery or Server-side discovery to determine the location of a service instance to which to send requests.

# MICROSERVICE REGISTRY / DISCOVERY

Issue

**How do clients of a service (in the case of Client-side discovery) and/or routers (in the case of Server-side discovery) know about the available instances of a service?**

# MICROSERVICE REGISTRY / DISCOVERY

## Forces

- Each instance of a service exposes a remote API such as HTTP/REST at a particular location (host:port)
- Number of services instances and their locations changes dynamically.
- Virtual machines and containers are assigned a dynamic IP address.
- Can adjust the number of instances based on load.

# MICROSERVICE REGISTRY / DISCOVERY

## Solution

- Implement a service registry, which is a database of services, their instances and their locations.
- Service instances are registered with the service registry on startup and deregistered on shutdown.
- Client of the service and/or routers query the service registry to find the available instances of a service.
- A service registry might invoke a service instance's health check API to verify that it is able to handle requests



## MICROSERVICE DISCOVERY (SERVER-SIDE)

# MICROSERVICE DISCOVERY (SERVER-SIDE)

## Issue

- How does the client of a service - the API gateway or another service - discover the location of a service instance?

# MICROSERVICE DISCOVERY (SERVER-SIDE)

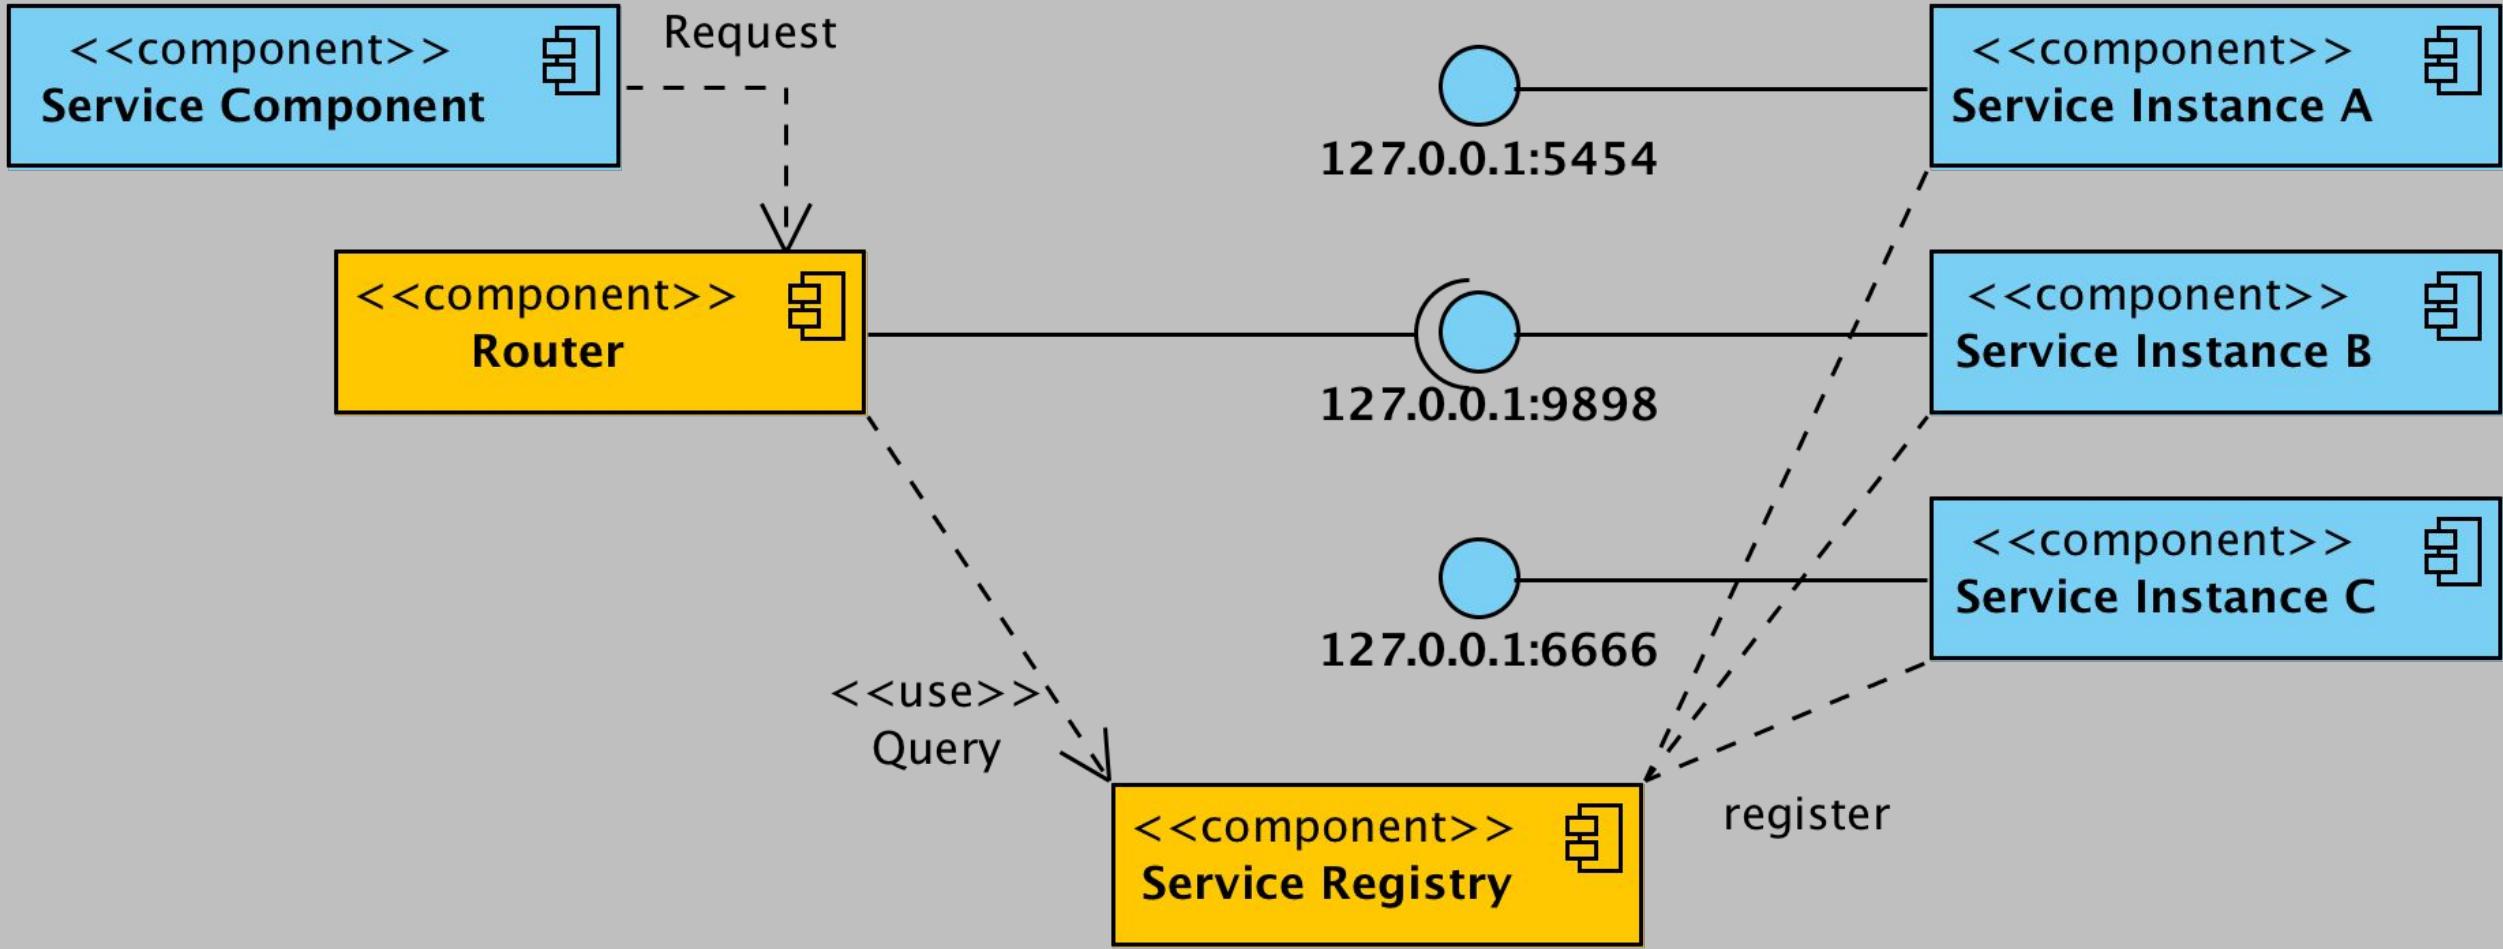
## Forces

- Each instance of a service exposes a remote API such as HTTP/REST at a particular location (host:port)
- Number of services instances and their locations changes dynamically.
- Virtual machines and containers are assigned a dynamic IP address.
- Can adjust the number of instances based on load.

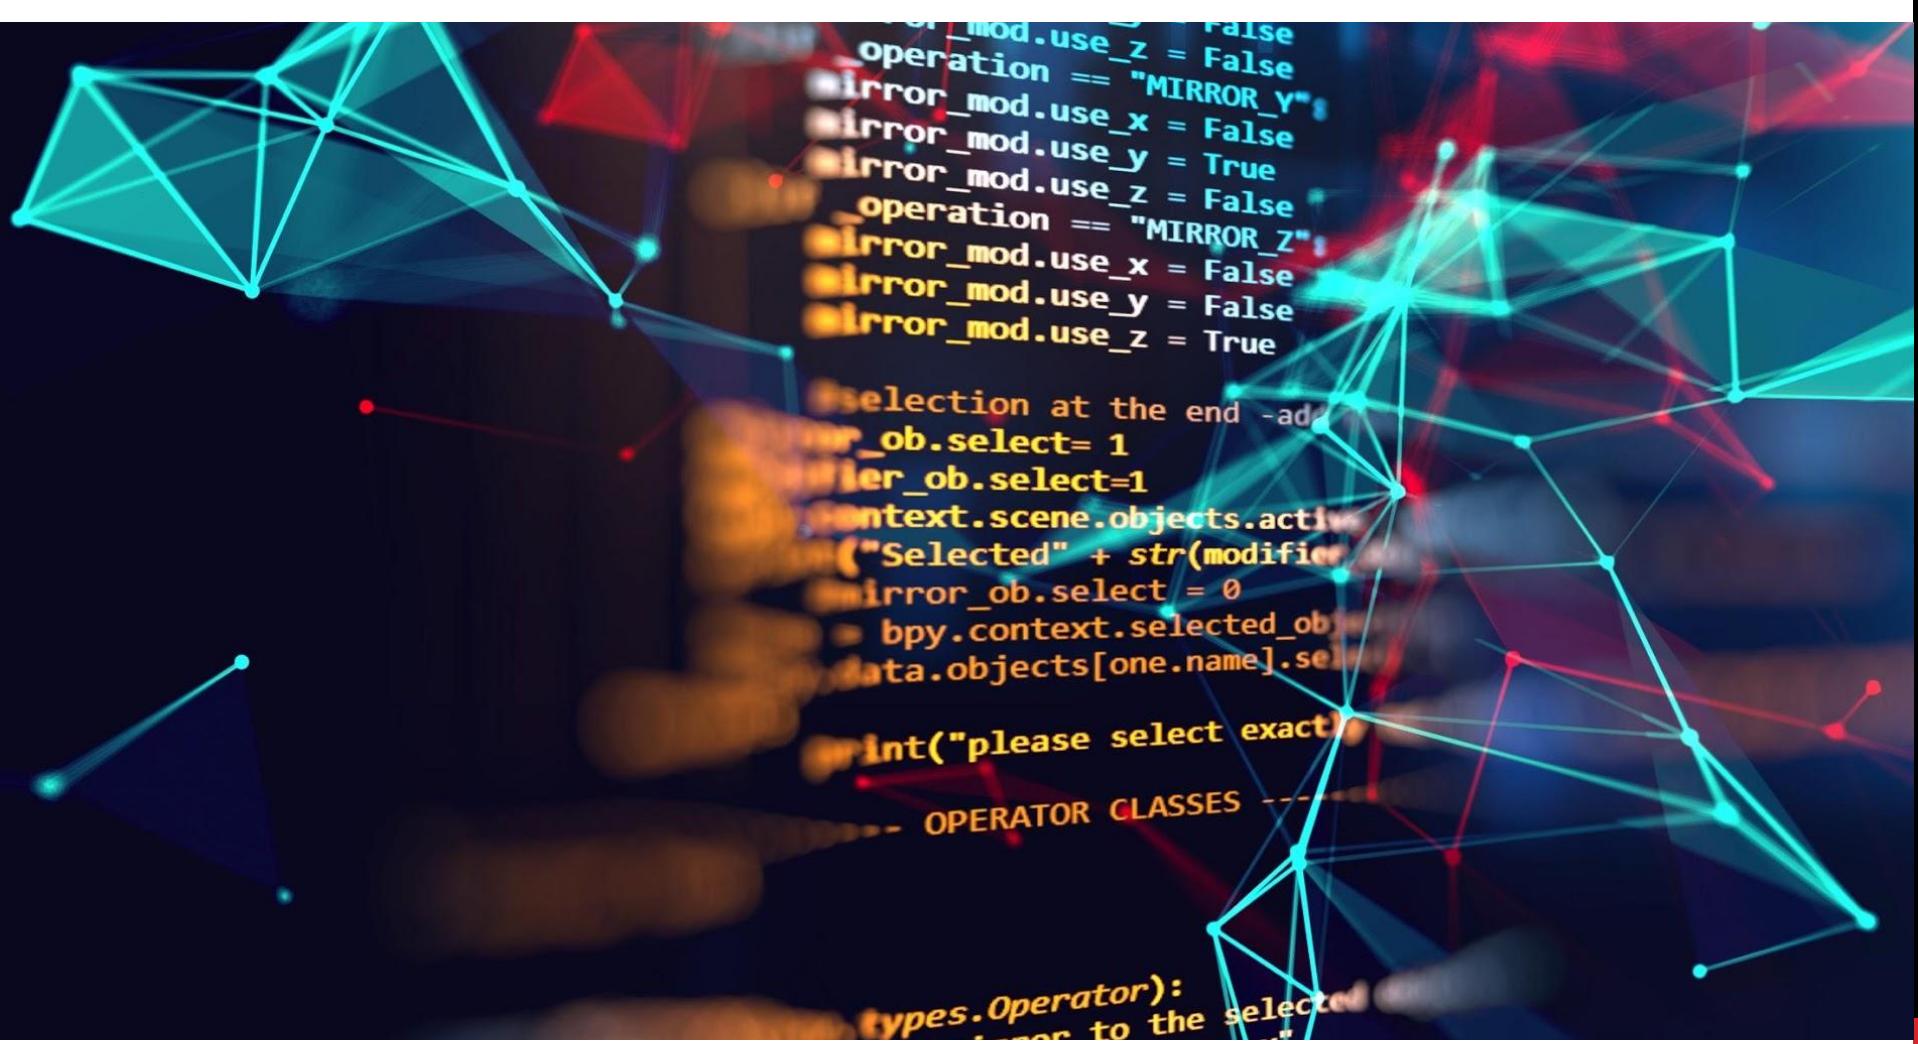
# MICROSERVICE DISCOVERY (SERVER-SIDE)

## Solution

- When making a request to a service, the client makes a request via a router (load balancer) that runs at a well known location.
- Router queries a service registry, which might be built into the router, and forwards the request to an available service instance.



## MICROSERVICE DISCOVERY (SERVER-SIDE)



# MICROSERVICE DISCOVERY (CLIENT-SIDE)

# MICROSERVICE DISCOVERY (SERVER-SIDE)

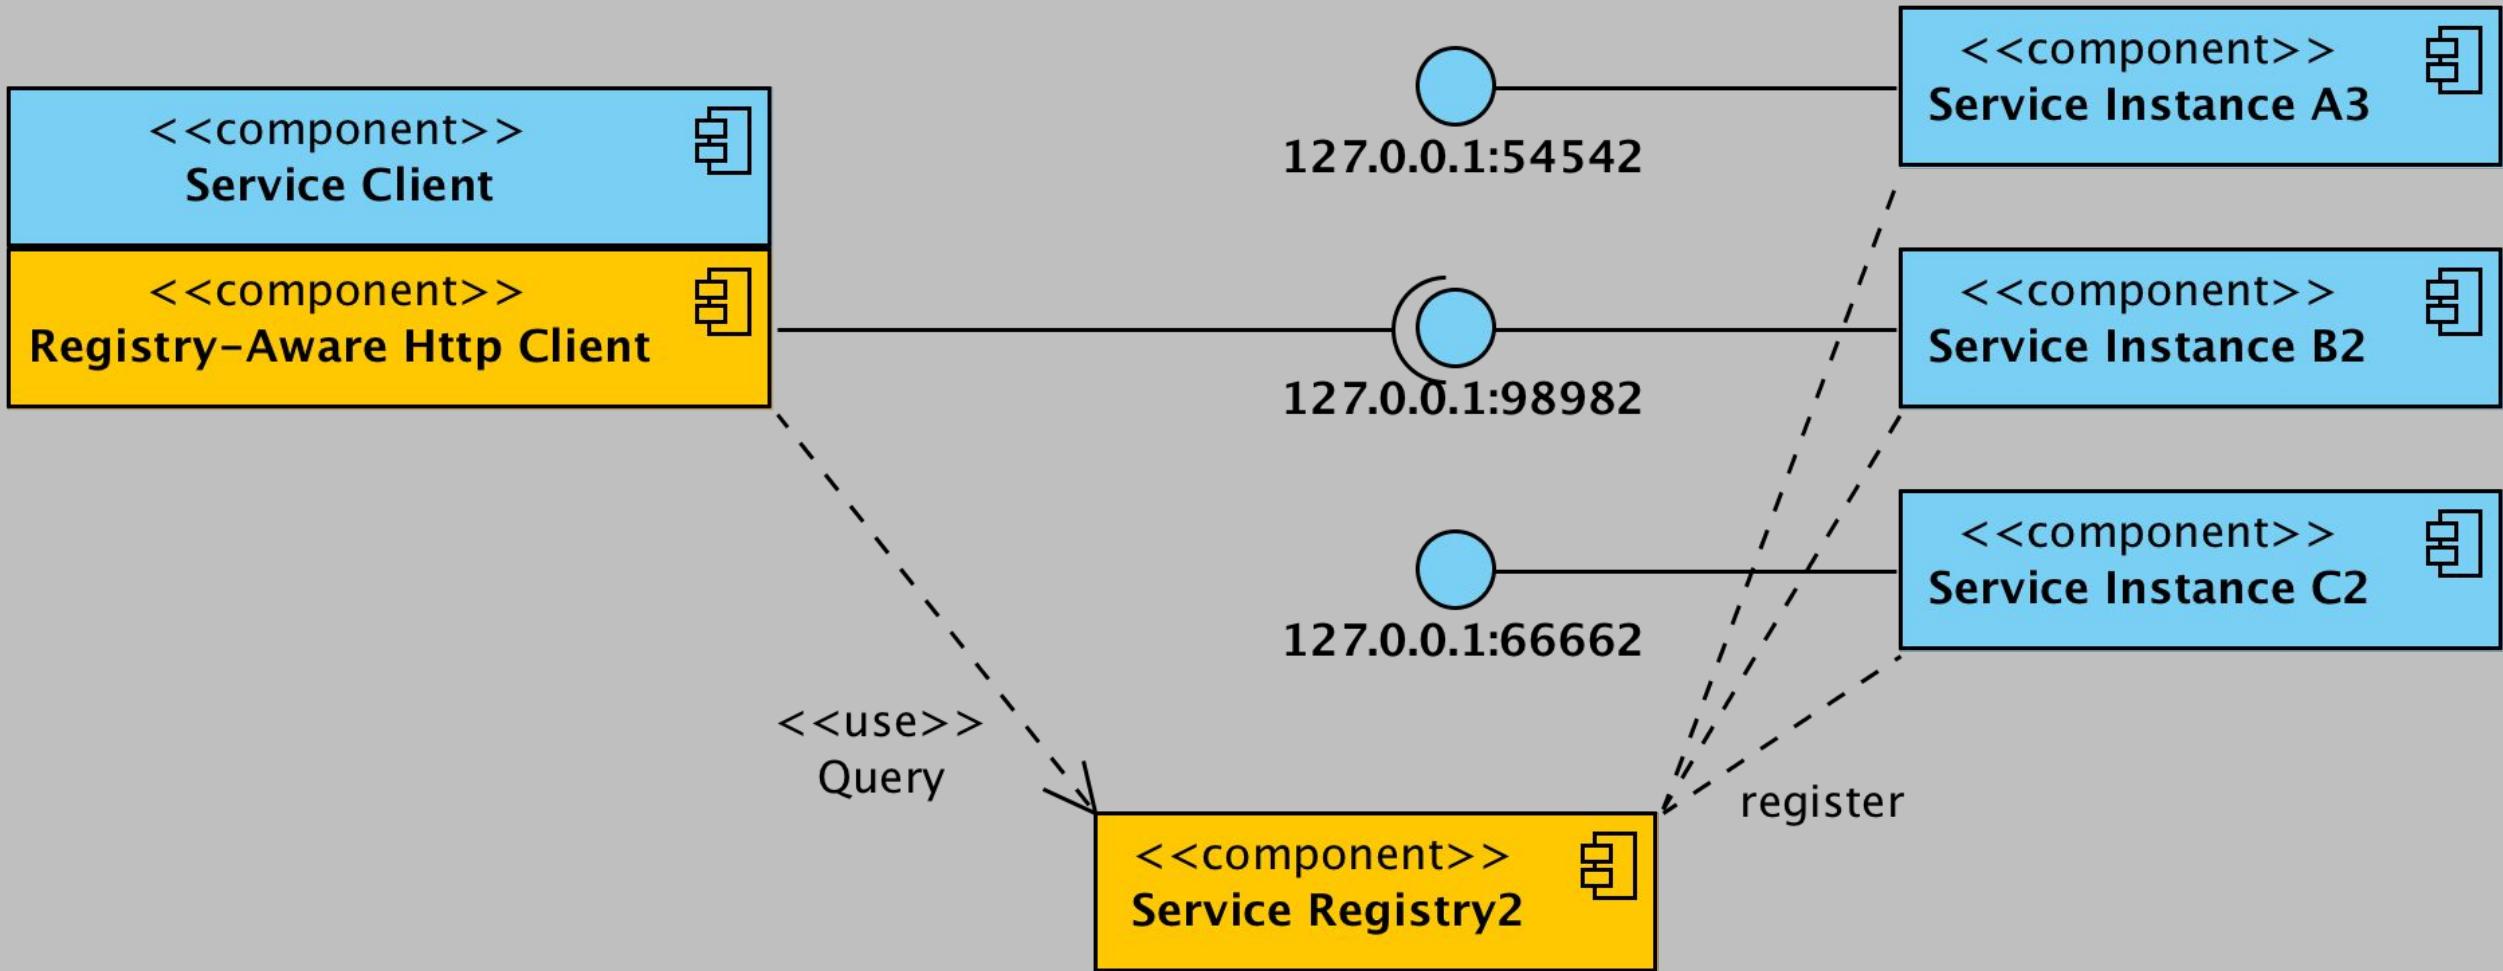
## Forces

- Each instance of a service exposes a remote API such as HTTP/REST at a particular location (host:port)
- Number of services instances and their locations changes dynamically.
- Virtual machines and containers are assigned a dynamic IP address.
- Can adjust the number of instances based on load.

# MICROSERVICE DISCOVERY (SERVER-SIDE)

## Solution

- When making a request to a service, the client obtains the location of a service instance by querying a Service Registry, which knows the locations of all service instances.

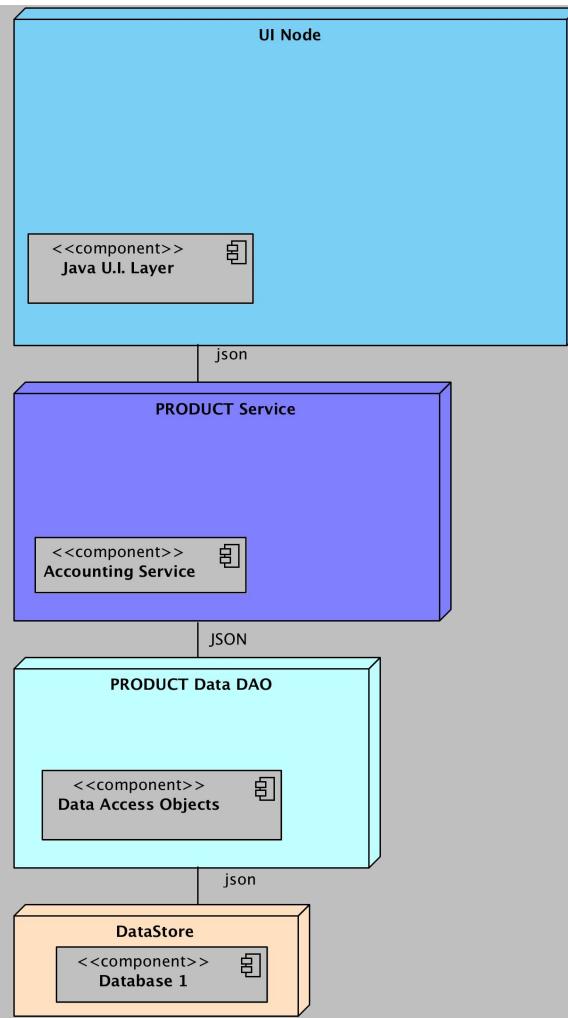


## MICROSERVICE DISCOVERY (CLIENT-SIDE)

# **FLEXIBLE SCALING OPTIONS**

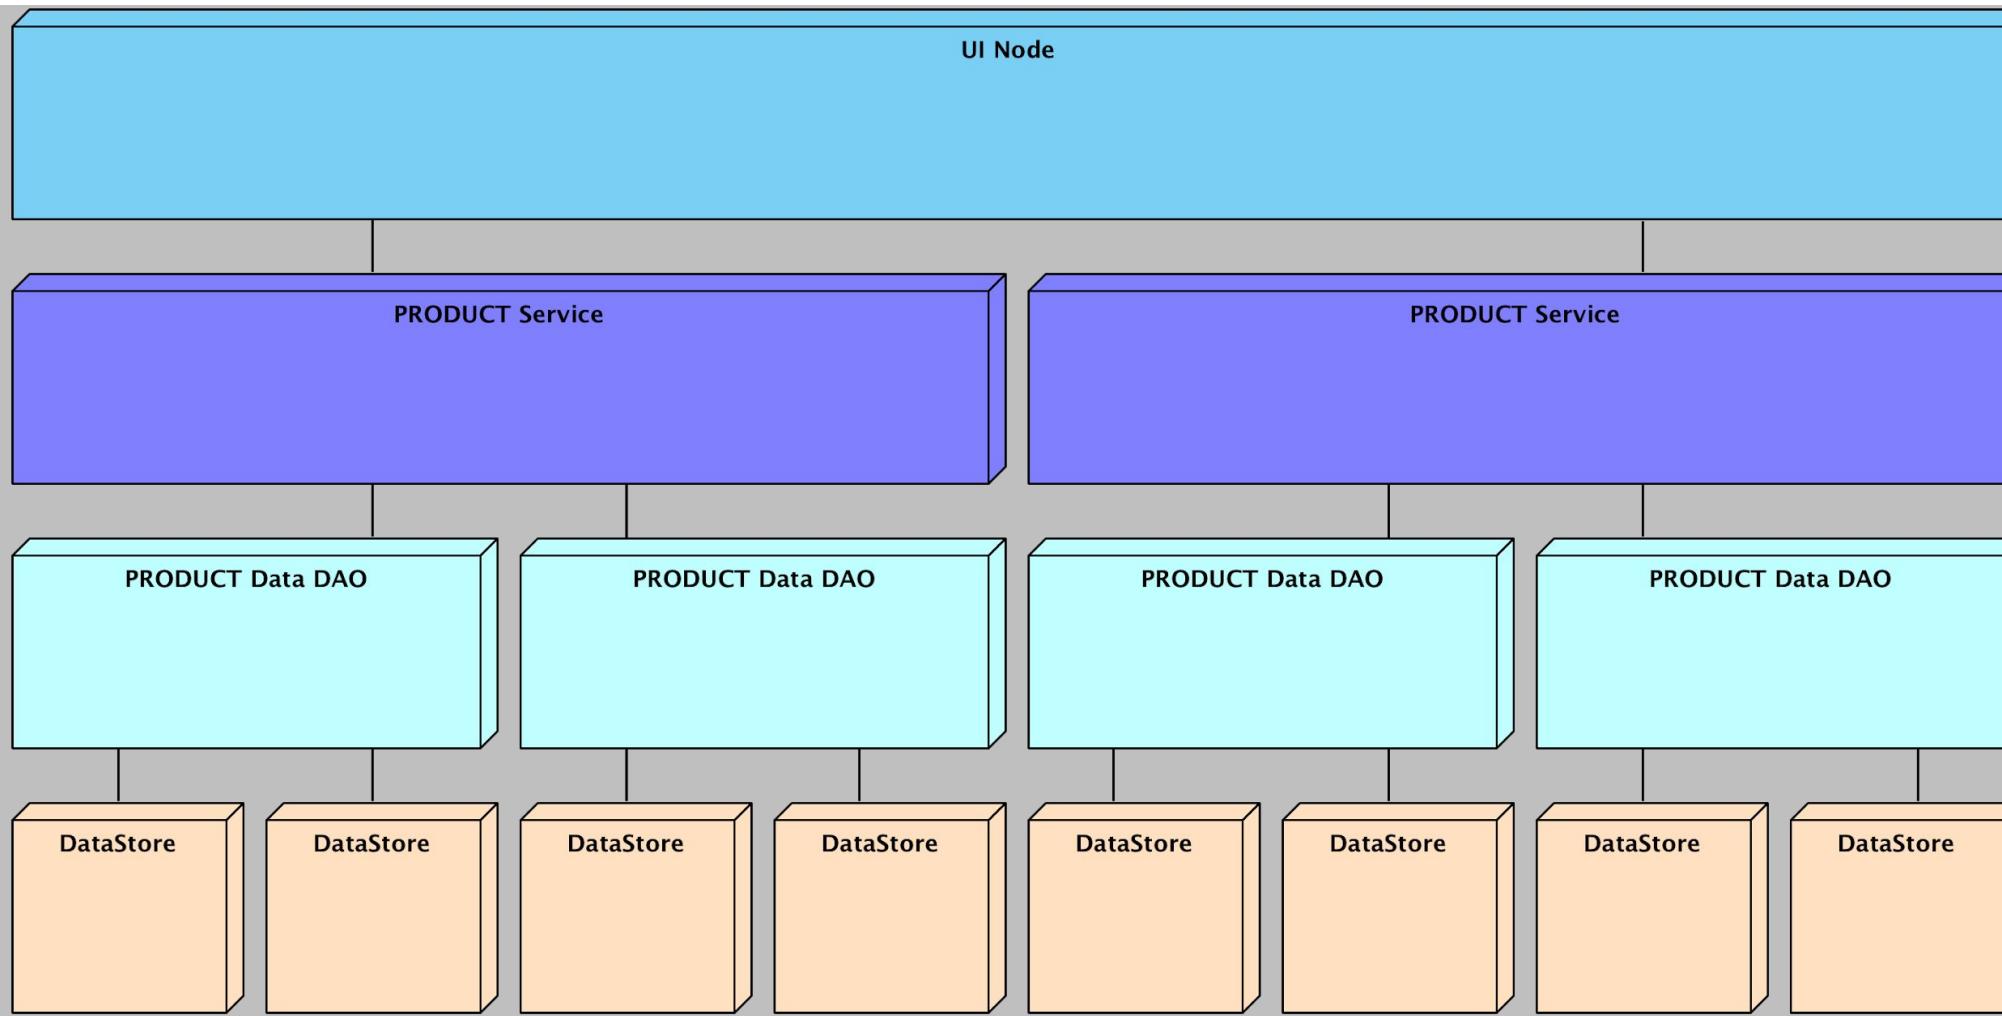
**Each service has:**

- Independent release schedule
- Independent load requirements
- Independent scaling strategy



## SCALING MICROSERVICES (VERTICALLY)

Larger instances



## SCALING MICROSERVICES (HORIZONTALLY)

More instances

# SERVICE ORCHESTRATION

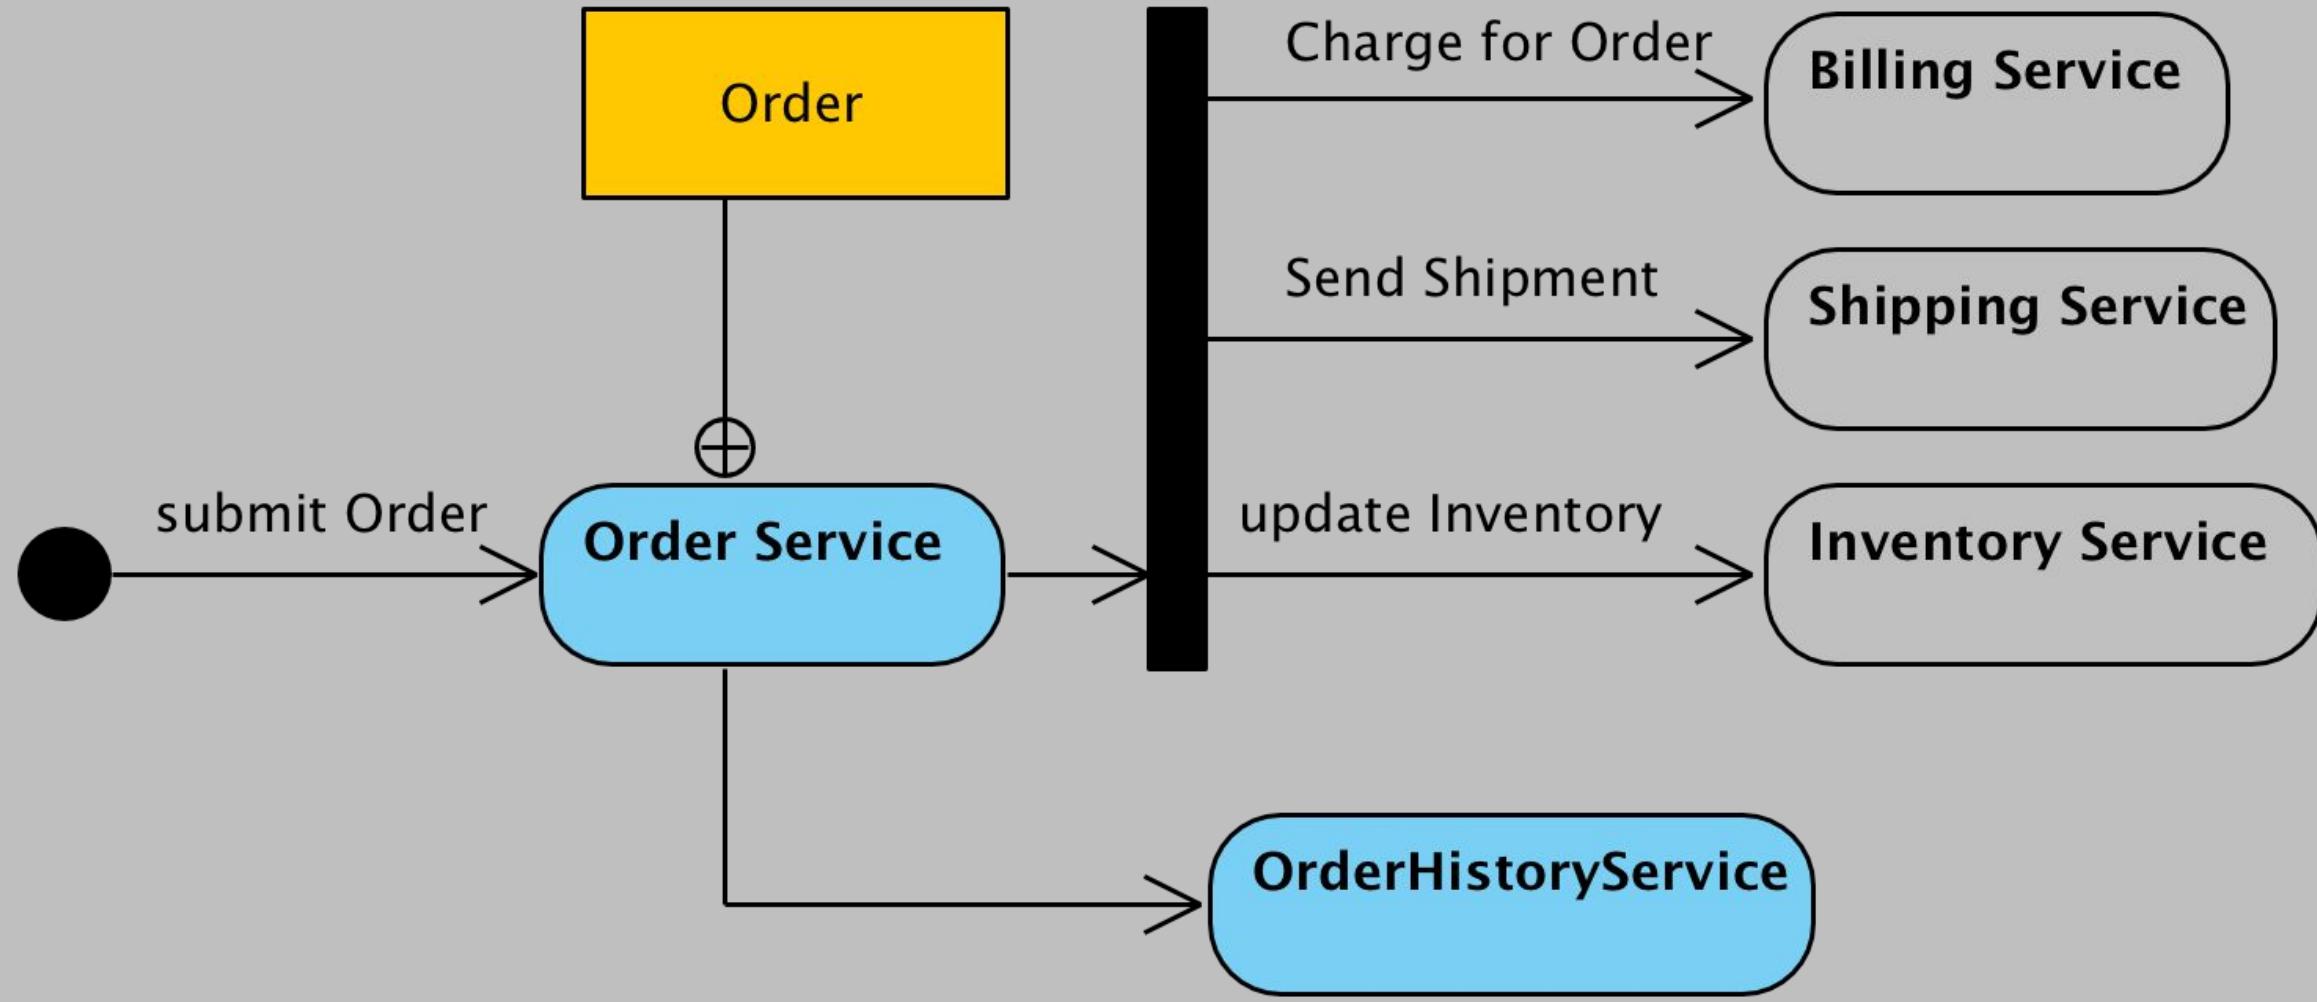
## Who is in charge of composing services

- In SOA, ESB often handles
- Service provider may handle
  - Orchestration uses a central service acts as the controller
    - Orchestrates the other services based on decision flow
  - Choreography has no central controller
    - A single event is published and sent to consumers
    - Each consumer of the event performs its function

## From microservices perspective, prefer choreography

- Controller as an external component or another microservice is counter to the autonomy principle

## There will be situations where some degree of orchestration is required



## SERVICE ORCHESTRATION

# ORCHESTRATION CASES

## Issue

**When orchestration is required, choices become more difficult and need to be examined carefully.**

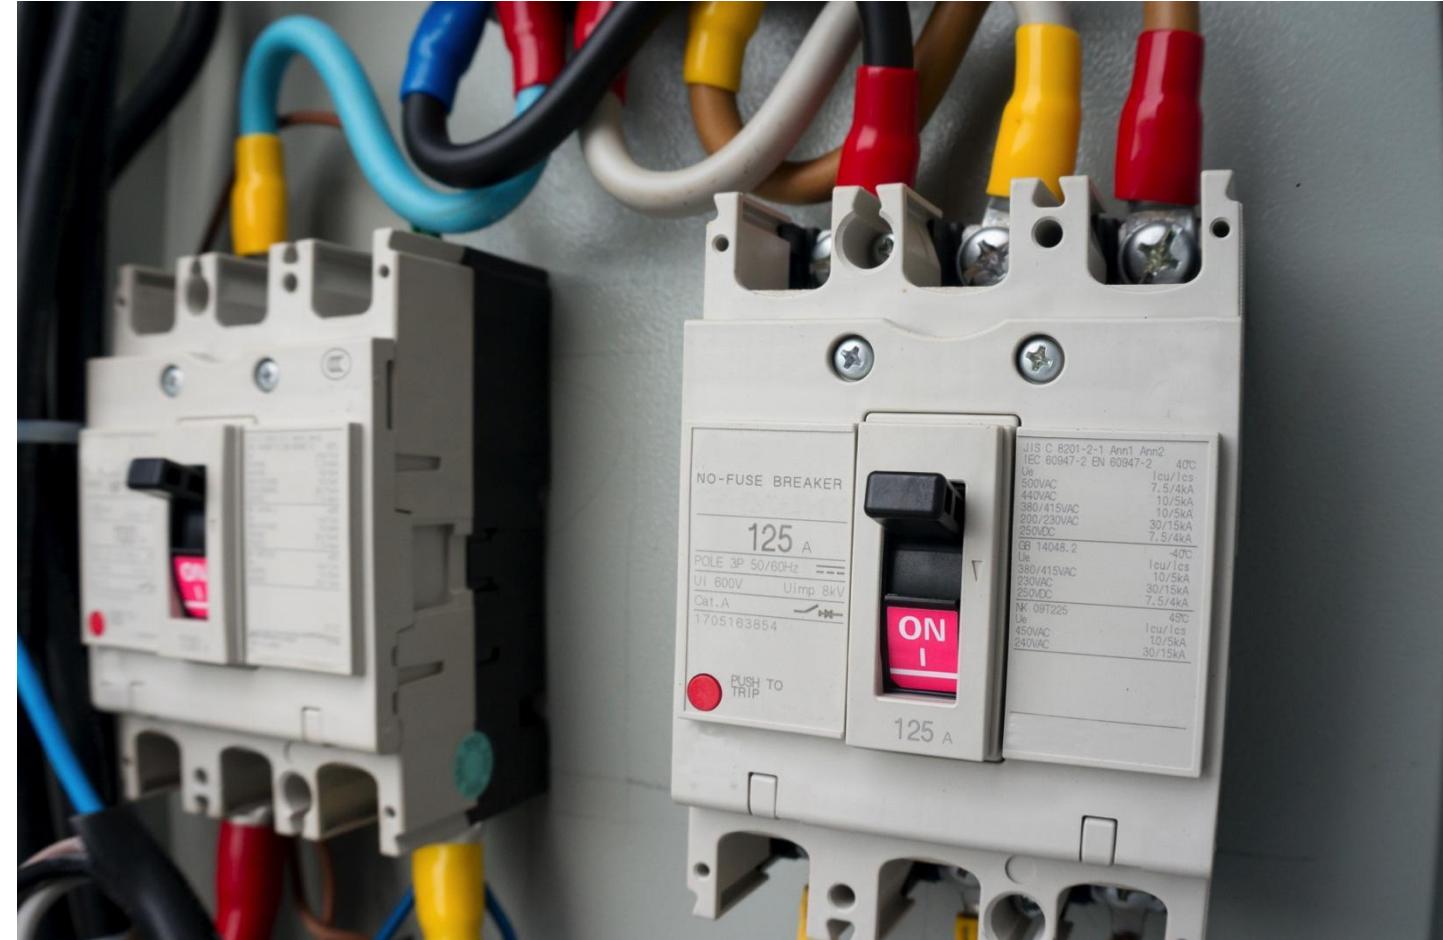
# ORCHESTRATION CASES

Options include:

- Adding in a synchronous call to each external service and wait for result
- Creating a specific microservice whose job it is to compose the component services
- Replicate the one or more of the component services into one of the higher level services

# CIRCUIT BREAKER

How to prevent a network or service failure from cascading to other services?



# CIRCUIT BREAKER

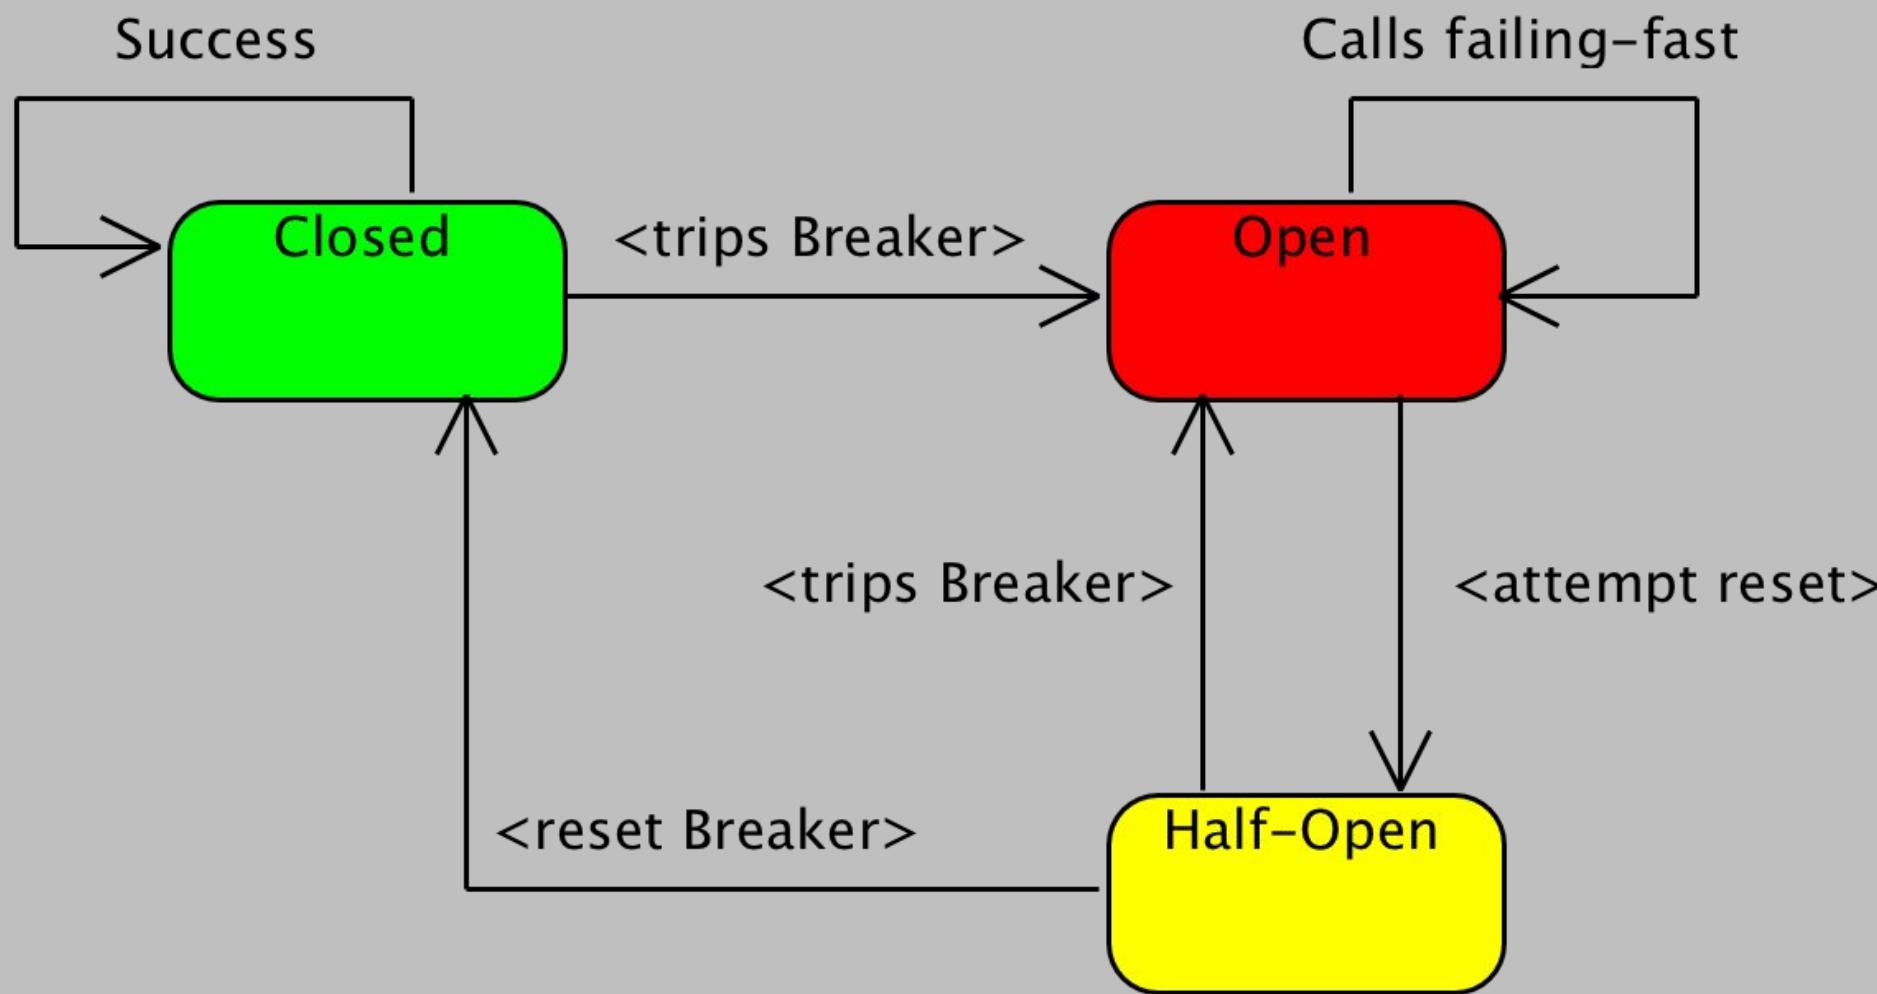
## Solution

- A service client should invoke a remote service via a proxy that functions in a similar fashion to an electrical circuit breaker.
- When the number of consecutive failures crosses a threshold, the circuit breaker trips, and for the duration of a timeout period all attempts to invoke the remote service will fail immediately.

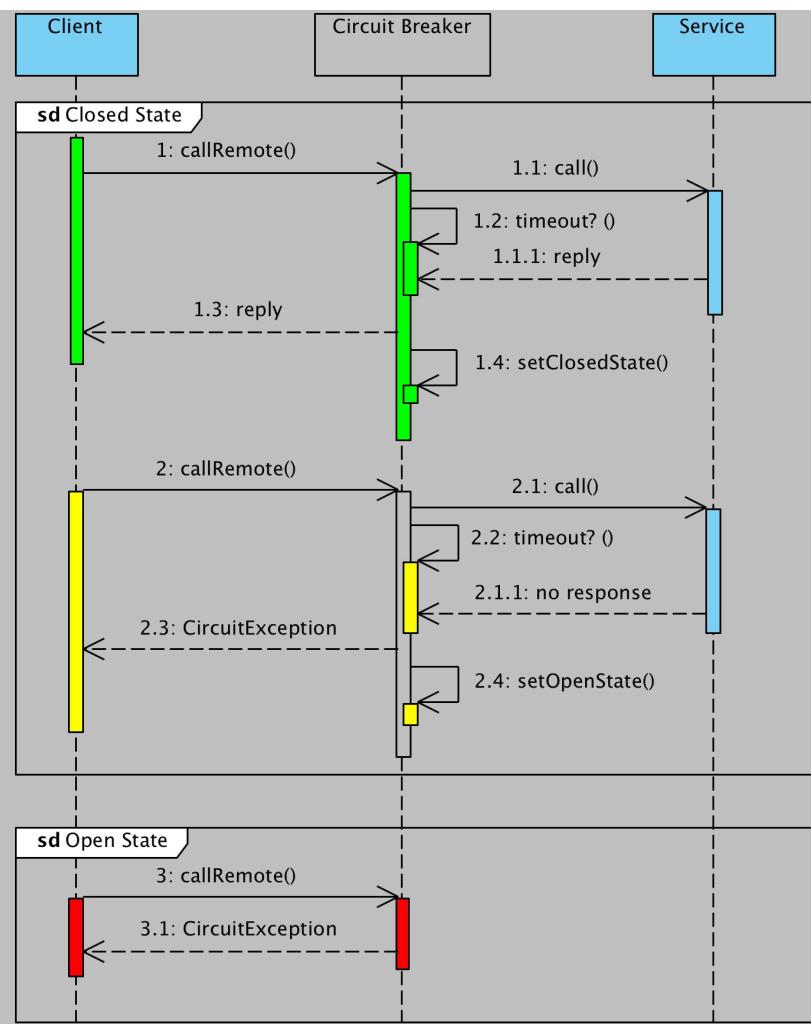
# CIRCUIT BREAKER (CONT'D)

## Solution

- After the timeout expires the circuit breaker allows a limited number of test requests to pass through.
- If those requests succeed the circuit breaker resumes normal operation.
- Otherwise, if there is a failure the timeout period begins again.



# CIRCUIT BREAKER



# CIRCUIT BREAKER

**THANK YOU  
FOR  
ATTENDING**